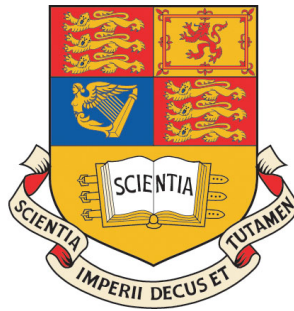


Density Matrix

Quantum Monte Carlo



Thomas William Rogers

Department of Physics

Imperial College London

MSci Project Report

Supervisors: Prof. Matthew Foulkes & Dr. James Spencer

Assessor: Dr. Alex Thom

Acknowledgements

Thank you to Prof. Matthew Foulkes and Dr. James Spencer who thought of the original idea for this project and provided a great deal of help both in the physics and the computing that resulted in the achievement of these results. Many thanks also to my project-partner Nicholas Blunt, with whom the work in this project is equally shared. Finally, I am grateful for discussions on the application of DMQMC in quantum information theory with Dr. Terry Rudolph and of Tom Poole for proof-reading this report and offering some important changes.

This work made use of the Imperial College High Performance Computing facilities.

Abstract

A new quantum Monte Carlo method based on stochastically sampling the thermal density matrix at finite temperatures has been developed. The method is an adaptation of the recently developed full configuration interaction quantum Monte Carlo[?] (FCIQMC) method and it aims to overcome two of FCIQMCs main restrictions by allowing the calculation of any quantum mechanical observable at finite temperatures. During this report the density matrix quantum Monte Carlo (DMQMC) algorithm is formulated by observing an analogy with FCIQMC. It is then tested on the spin-1/2 antiferromagnetic Heisenberg model on a bipartite lattice.

The DMQMC method produces results that agree with the exact finite temperature FCI energy and exact ground-state staggered magnetisation for the 4×4 lattice. For the larger 6×6 lattice the standard DMQMC algorithm fails. At this point an importance sampling method is introduced, yielding results comparable to accurate Green's function quantum Monte Carlo (GFQMC) estimates. The importance sampling method also enabled simulation of the $4 \times 4 \times 4$ lattice, producing results for the ground-state energy that matched those obtained by FCIQMC.

Finally an avenue into quantum information theory is explored and some key results for the entanglement between two qubits on a spin-1/2 antiferromagnetic Heisenberg ring are reproduced. Additionally, both how the entanglement between qubits changes in a uniformly applied magnetic field and with increasing separation between the two qubits are investigated.

List of Figures

Chapter 1

Introduction

The main aim of current electronic structure methods is to calculate the ground-state properties of many-electron systems. Conventional full-configuration interaction (FCI) calculations quickly become unfeasible because the size of the Hilbert space grows exponentially with the number of electrons. Such calculations based on iterative diagonalisation require the storage of at least as many double-precision numbers as twice the size of the Hilbert space[?].

Quantum Monte Carlo (QMC) methods can provide accurate simulations of many-electron systems via a stochastic sampling of the many-electron wave function. Such methods are favourable if the number of agents or “walkers” conducting the stochastic sampling is a small fraction of the size of the Hilbert space.

Projector methods[?] such as diffusion Monte Carlo (DMC) rely on the ground-state emerging from the imaginary-time Schrödinger equation in the limit of infinite imaginary-time (see Appx. A). However, because DMC does not enforce the constraint that a Fermionic wave function must be anti-symmetric, they suffer from convergence to the incorrect bosonic ground-state in an event known as the “bosonic catastrophe”. In such cases the fixed-node approximation is applied[?] [?], where *a priori* knowledge of the wave-function is used to restrict the propagation of walkers so that they evolve towards an approximately antisymmetric solution.

Until recently QMC methods, such as DMC, could only provide numerically exact solutions for systems of many interacting bosons. The “fermion sign problem”, described above, prevented an exact solution for systems of many interact-

ing fermions.

However, recently Booth *et al.*[?] introduced a new approach, known as full configuration-interaction QMC (FCIQMC), in which the antisymmetric property of the many-electron wave function is enforced by working in a discrete space of slater determinants (see Appx. A). The key is the improved efficiency of walker annihilation, which ensures convergence to the fermionic ground-state solution[?].

FCIQMC has two main restrictions. The first is that it proves difficult to calculate the expectation values of operators that do not commute with the Hamiltonian. The second being that it only allows for calculation of expectation values in the ground-state. This project introduces a method that is similar to FCIQMC, but attempts to overcome both of these restrictions. This is done by a stochastic sampling of the many-electron thermal density matrix as a function of inverse-temperature, rather than sampling the many-electron wave function as a function of imaginary-time. This in theory, allows for the calculation of the expectation values of any operator at finite temperatures.

In Ch. ?? the density matrix quantum Monte Carlo (DMQMC) algorithm is developed and an optimised and highly parallel implementation is discussed. In Ch. ?? DMQMC is applied to the spin-1/2 antiferromagnetic Heisenberg model on two and three dimensional bipartite lattices. The results are compared against exact values from the full diagonalisation of the Hamiltonian, for small systems, and accurate FCIQMC or Green's function QMC (GFQMC) estimates for larger systems. Finally, in Ch. ?? a method of calculating a reduced density matrix for a subsystem of the Heisenberg lattice is developed and this leads into applications in quantum information theory.

Note that a system of units with $\hbar = k_B = 1$ is adopted for all equations in this report.

Chapter 2

The DMQMC algorithm

In this chapter the DMQMC algorithm is formulated using an analogy between the imaginary-time Schrödinger equation and the evolution of the many-electron thermal density matrix as a function of inverse-temperature. Several features of FCIQMC are used in the construction of this algorithm in a way that maintains FCI-quality results. Finally, the main features of an optimised and highly parallel implementation of the DMQMC algorithm are discussed.

2.1 The thermal density matrix

In quantum mechanics the density operator[?] is any operator, ρ , that satisfies three properties:

1. It is Hermitian i.e. $\rho = \rho^\dagger$
2. It is a positive, semi-definite operator, i.e. $\langle \psi | \rho | \psi \rangle \geq 0 \quad \forall \quad |\psi\rangle \in \mathcal{H}$
3. It has unit trace i.e. $\text{Tr}\rho = 1$

As ρ is Hermitian, it can be expanded using the spectral theorem,

$$\rho = \sum_{j=1} p_j |\psi_j\rangle \langle \psi_j|, \quad (2.1)$$

where $\{p_j\}$ are the eigenvalues of ρ and are also the classical probabilities that the system is in the vector state $|\psi_j\rangle$. A density operator is useful for describing

and calculating the properties of a mixed state. A mixed state is a statistical ensemble of several quantum states and arises in situations where there is classical uncertainty (which is different to quantum uncertainty).

One major benefit of constructing a density matrix is that it allows for very simple calculations of any quantum mechanical observable, O :

$$\begin{aligned}
\langle O \rangle &= \sum_i p_i \langle \psi_i | O | \psi_i \rangle \\
&= \sum_i \langle \psi_i | \rho | \psi_i \rangle \langle \psi_i | O | \psi_i \rangle \\
&= \sum_i \langle \psi_i | \rho O | \psi_i \rangle \\
&= \sum_{i,j} \rho_{ij} O_{ji}.
\end{aligned} \tag{2.2}$$

In the last line the matrix representations of the operators ρ and O have been adopted.

This hints that the formulation of a QMC method that can stochastically sample the many-electron density operator, could provide a simple calculation of the expectation value of any quantum mechanical observable. There is no restriction on the operator, O , other than the requirement that it is expressed in the same basis as the density matrix. Eq. ?? is the blue-print used for calculating all estimators of quantum mechanical observables in this section.

An important mixed state is the thermal state[?], which describes a quantum statistical system at temperature T . In this case, the classical probability distribution is described by the Boltzmann distribution and the thermal density operator is

$$\rho = \frac{e^{-\beta H}}{Z(\beta)}, \tag{2.3}$$

where H is the Hamiltonian operator, $\beta = 1/T$ is the inverse-temperature in natural units and $Z(\beta) = \text{Tr}(e^{-\beta H})$ is the partition function of the system.

Applying the spectral theorem to the Hamiltonian, Eq. ?? becomes

$$\rho = \sum_j \frac{e^{-\beta E_j}}{Z(\beta)} |E_j\rangle \langle E_j|. \tag{2.4}$$

This reflects the classical uncertainty induced by thermal fluctuations of the energy state that the system is in. The canonical partition function can be expressed as

$$Z(\beta) = \sum_{j=1} e^{-\beta E_j}. \quad (2.5)$$

From now on the convention

$$\rho = e^{-\beta H} \implies \langle O \rangle = \frac{\sum_{i,j} \rho_{ij} O_{ji}}{\sum_i \rho_{ii}} \quad (2.6)$$

is adopted to keep notation simple.

2.2 An analogy with FCIQMC

Using the convention from Eq. ?? it is evident that, in its matrix representation, ρ obeys the symmetrised Bloch equation^{? i}

$$\frac{\partial \rho_{ij}}{\partial \beta} = -\frac{1}{2} \sum_k H_{ik} \rho_{kj} + \rho_{ik} H_{kj}, \quad (2.7)$$

with the initial condition

$$\boldsymbol{\rho}(\beta = 0) = \boldsymbol{I}, \quad (2.8)$$

where \boldsymbol{I} is the identity matrix with the same dimensions as the thermal density matrix $\boldsymbol{\rho}$.

Now as the partition function is simply $\text{Tr} \boldsymbol{\rho}$, Eq. ?? provides all of the information required to calculate the normalised thermal density matrix. The Bloch equation (Eq. ??) is analogous to the imaginary-time Schrödinger equation. The columns of the many-electron density matrix evolving as a function of inverse-temperature, rather than a many-electron wave function evolving in imaginary-time,

$$\frac{\partial \psi_i}{\partial \tau} = -\sum_j H_{ij} \psi_j, \quad (2.9)$$

ⁱThe symmetrised version of the Bloch equation is used so that there is spawning in two directions rather than one, which makes more sense intuitively.

where ψ_i is the vector representation of the many-electron wave function in a discrete space (e.g. Slater determinant space, see Appx. A) and $\tau = it$ is the imaginary-time.

In the true spirit of this analogy, the ground-state many-electron density matrix is projected out in the large β (or zero temperature) limit,

$$\rho(\beta \rightarrow \infty) = \lim_{\beta \rightarrow \infty} \sum_i e^{-\beta E_i} |E_i\rangle \langle E_i| \approx e^{-\beta E_0} |E_0\rangle \langle E_0|, \quad (2.10)$$

where $|E_0\rangle$ is the ground-state or lowest eigenvector of H and E_0 is the ground-state energy eigenvalue.

This is similar to in FCIQMC, where the ground-state eigenvector, ψ_0 , of the many-electron Hamiltonian matrix, \mathbf{H} , is projected out in the long imaginary-time limit[?],

$$\begin{aligned} \psi(\tau \rightarrow \infty) &= \lim_{\tau \rightarrow \infty} e^{-\mathbf{H}\tau} \sum_{\alpha} v_{\alpha}(0) \psi_{\alpha} \\ &= \lim_{\tau \rightarrow \infty} \sum_{\alpha} v_{\alpha}(0) e^{-E_{\alpha}\tau} \psi_{\alpha} \\ &\approx v_0(0) e^{-E_0\tau} \psi_0, \end{aligned} \quad (2.11)$$

where the initial vector has been expanded in terms of the complete orthonormal set of eigenvectors, $\{\psi_{\alpha}\}$ of \mathbf{H} i.e. $\psi(0) = \sum_{\alpha} v_{\alpha}(0) \psi_{\alpha}$.

So a QMC method based on solving Eq. ?? would evolve towards the ground-state of the many-electron system in a similar way to in FCIQMC, however it would also allow for the density matrix to be sampled at finite temperatures.

The difficult factor of $e^{-\beta E_0}$ in Eq. ?? can be removed by setting $E_0 = 0$. In reality E_0 is usually unknown and so instead one can solve,

$$\frac{\partial \rho_{ij}}{\partial \beta} = \frac{1}{2} \sum_k T_{ik} \rho_{kj} + \rho_{ik} T_{kj}, \quad (2.12)$$

where $T_{ik} = -(H_{ik} - S\delta_{ik})$ is the “update matrix” and S is the energy shift, which can be adjusted carefully to keep normalisation approximately constant[?].

2.3 Population dynamics

So drawing upon the analogy between Eq. ?? and Eq. ??, FCIQMC can be adapted so that rather than stochastically sampling the many-electron wave function, it performs a stochastic sampling of the many-electron thermal density matrix. As in FCIQMC, one considers a collection of markers or “psips”ⁱ that are distributed over the elements of the density matrix. Now each psip still has a “charge” $q = \pm 1$, but their location is now labelled by two indices (i, j) .

The fact that the solution to Eq. ?? is a matrix implies that the psips can now spawn from two ends. One end can spawn in one direction and the other end can spawn in a direction that is perpendicular. The algorithm for the DMQMC can be summarised as follows:

At each inverse-temperature step $\Delta\beta$, loop over the population of psips and allow each end, A and B , of the psip to spawn a “child” at other locations according to the following set of rules:

1. The probability that end A of a psip at (i, j) spawning a child at (k, j) is $\frac{1}{2} |T_{ki}| \Delta\beta$
2. The probability that end B of a psip at (i, j) spawning a child at (i, k) is $\frac{1}{2} |T_{kj}| \Delta\beta$
3. If end A of a parent psip with charge q_{parent} at location (i, j) spawns a child at (k, j) , the charge of the child is given by $q_{child|A} = \text{sign}(T_{ki})q_{parent}$
4. If end B of a parent psip with charge q_{parent} at location (i, j) spawns a child at (i, k) , the charge of the child is given by $q_{child|B} = \text{sign}(T_{kj})q_{parent}$

At the end of each inverse-temperature step, after each psip has spawned as many times as it can, pairs of walkers of opposite charge at the same location annihilate each other and are removed from the simulation.

ⁱ“psips” stands for “psi particles”[?] but there was a temptation to call them “rhps”

2.4 First order finite-difference approximation

The motivation for this algorithm is that the dynamics satisfy a first-order Euler finite-difference approximation of Eq. ???. The expected charge, $\bar{q}_{ij}(\beta + \Delta\beta)$, at location (i, j) at inverse-temperature $\beta + \Delta\beta$ is related to the expected charges $\bar{q}_{kj}(\beta)$ and $\bar{q}_{ik}(\beta)$ at inverse-temperature β by

$$\bar{q}_{ij}(\beta + \Delta\beta) = \bar{q}_{ij}(\beta) + \frac{1}{2} \sum_k (T_{ik} \bar{q}_{kj} + \bar{q}_{ik} T_{kj}) \Delta\beta. \quad (2.13)$$

The first term on the right-hand side describes the total charge of the walkers that were at (i, j) at inverse-temperature β . The second term describes the spawning of walkers onto (i, j) over the inverse-temperature interval $\Delta\beta$. As this is a first-order Euler finite approximation of Eq. ??, the distribution of charges is proportional to the density matrix.

The stability[?] of the first order finite-difference approximation in Eq. ?? can be analysed by making a linear perturbation to the density matrix i.e $\rho \rightarrow \rho + \epsilon$. This implies that for the n^{th} β -step,

$$\epsilon_{n+1} = (\mathbf{I} + \Delta\beta \mathbf{T}) \epsilon_n. \quad (2.14)$$

Now this is stable if $|\lambda^i| \leq 1$, where $\{\lambda^i\}$ are the eigenvalues of $(\mathbf{I} + \Delta\beta \mathbf{T})$. So this leads to the stability condition, $\Delta\beta \leq 2/(E_{max} - S)$, where E_{max} is the maximum energy eigenvalue of the Hamiltonian. The shift, S , can be assumed to be approximately equal to the energy of the system and is most negative at the ground-state. So in the worst case $\Delta\beta$ is limited to

$$0 < \Delta\beta \leq \frac{2}{E_{max} - E_0}. \quad (2.15)$$

Furthermore, by considering the Taylor expansion of the density matrix at inverse-temperature $\beta + \Delta\beta$, it is evident that the local error for a density matrix element is $\mathcal{O}(\Delta\beta^2)$. Note that statistical errors will also be accrued because Eq. ?? is integrated stochastically. This analysis of the accuracy and the stability was taken into consideration when selecting the $\Delta\beta$ in order to avoid instability or unnecessarily large errors. For all simulations that follow in this report $\Delta\beta = 0.1$

was chosen for stability and also as a compromise between accuracy and efficiency.

2.5 Shift update algorithm

The algorithm for updating the shift S was first used by Umrigar[?] in DMC but was also adopted by the inventors of FCIQMC[?]. In this, S is adjusted according to

$$S(\beta + A\Delta\beta) = S(\beta) + \frac{\zeta}{A\Delta\beta} \ln \left(\frac{N_p(\beta + A\Delta\beta)}{N_p(\beta)} \right), \quad (2.16)$$

where A is the number of β -steps between shift updates, ζ is a shift damping parameter and $N_p(\beta)$ is the total number of psips at the inverse-temperature β . During simulations, ζ is chosen carefully to prevent large fluctuations in S . It should also be noted that by averaging the shift over many iterations, it can be used as an estimator for the ground-state energy in a similar way to in FCIQMC. However, this approach was not adopted during this project.

2.6 The spin-1/2 Heisenberg model

In 1926, shortly after the birth of quantum mechanics, Werner Heisenberg[?] and Paul Dirac[?] independently recognised that the new theory could be used to describe the mysterious properties of ferromagnetism. They realised that quantum mechanics implied the existence of an effective interaction, the exchange interaction, between electron spins that is caused by the combined effects of the Pauli exclusion principle and the Coulomb repulsion between them[?].

The Hamiltonian that describes the exchange interactions between a set of quantum mechanical spins (with periodic boundary conditions) is

$$H = -J \sum_{\langle i,j \rangle} S_i S_j + h \sum_i S_i^z, \quad (2.17)$$

where $S = \{S^x, S^y, S^z\} = \frac{1}{2}\{\sigma^x, \sigma^y, \sigma^z\}$ are the quantum spin operators expressed in terms of the Pauli spin matrices, h is a uniform magnetic field in the z -direction and J is the coupling constant. Here it has been assumed that nearest-neighbour interactions are dominant, hence the summation is only over nearest-neighbours.

Furthermore, it is assumed that the system is isotropic so that J is the same between any two neighbouring spins.

The coupling constant J plays a big role in determining the physical properties of the system. If there is no externally applied magnetic field ($h = 0$) then for $J > 0$ the ground-state is always ferromagnetic, where neighbouring spins tend to align themselves in the same direction. On the other hand, for $J < 0$, the ground-state is antiferromagnetic, where neighbouring spins point in opposite directions.

For the Heisenberg model with N spins, H acts on a 2^N dimensional Hilbert space, which is spanned by the orthogonal basis vectors, or spin configurations, $|s_1 s_2 \dots s_N\rangle$. Here, s_n can represent either a spin in the positive z -direction (spin up), $s_n = \uparrow$, or a spin in the negative z -direction (spin down), $s_n = \downarrow$. On implementing DMQMC for the Heisenberg model, bitstrings can be used to represent a given spin-configuration of the system. For example, the configuration for a system of N spins can be represented by a string of N binary digits, with a 1 representing a spin up and 0 a spin down.

The Heisenberg Hamiltonian, Eq. ?? can also be expressed in terms of the spin flip operators $S^\pm = S^x \pm iS^y$,

$$H = -J \sum_{\langle i,j \rangle} \left[\frac{1}{2} (S_i^+ S_j^- + S_i^- S_j^+) + S_i^z S_j^z \right] + h \sum_i S_i^z. \quad (2.18)$$

This gives a more intuitive feel for the action of the Hamiltonian on a given configuration in the Heisenberg model. Every non-zero off-diagonal element of the Hamiltonian flips a pair of neighbouring spins from $\uparrow\downarrow$ to $\downarrow\uparrow$ and vice versa. So a matrix element, H_{ij} , has a non-zero value of $-J/2$, if the configurations $|i\rangle$ and $|j\rangle$ differ by a spin-flip of neighbouring anti-aligned spins. In other words, the two configurations are a single excitation apart.

For the diagonal elements, H_{ii} , neighbouring spins in the configuration, $|i\rangle$, that are aligned ($\uparrow\uparrow$ or $\downarrow\downarrow$) give a contribution of $-J/4$. Neighbouring spins that are anti-aligned ($\downarrow\uparrow$ or $\uparrow\downarrow$) give a contribution of $J/4$. For a single total spin (M_S) subspaceⁱ, the magnetic field only affects the diagonal elements of the

ⁱThe current DMQMC implementation only considers a single M_S subspace during a simulation. This is because the code is built upon an FCIQMC code, which only requires simulations

Hamiltonian,

$$\mathbf{H}(h) = \mathbf{H}(h = 0) + \frac{hM_S}{2}\mathbf{I}. \quad (2.19)$$

2.7 Implementation

The DMQMC algorithm was implemented in `FORTRAN 90` (see Appdx. C) and was an adaptation of James Spencer’s optimised and highly parallel FCIQMC code. As this code adopted the same spawning and death/clone steps as the original FCIQMC method[?], the DMQMC algorithm has been implemented in a slightly different way to in Sec. ???. To be specific, the effect of the shift is simulated by a death/clone step rather than by changing the spawning probabilities.

In a simulation, the statistics for each β -step are collected from a number of different “ β -loops”, each of which start with a different random number generator (RNG) seed. In each β -loop the main steps are:

1. Initiate the psip population on the diagonal of the density matrix
2. Update DMQMC estimators and initialise shift
3. Move to next β -step
4. Loop over all psips and attempt to spawn progeny onto two other connected density matrix elements
5. At the same time, attempt to kill or clone each parent psipⁱ
6. Annihilate psips with opposite charges that occupy the same density matrix element
7. Update DMQMC estimators and shift (if necessary)
8. Repeat steps 3-7 until the user-defined maximum β has been reached

in a single M_S subspace, the one that contains the ground-state. Unfortunately, there was not enough time during this project to combine all M_S subspaces.

ⁱThe combination of steps 4 and 5 is equivalent to the spawning step described in Sec. ?? using the update matrix T . This was how the FCIQMC algorithm was originally proposed by Booth *et al.*[?].

In this implementation the position of a psip within the density matrix is represented by two bitstrings, with one corresponding to row position and the other to column position. The bitstring can also be seen as a representation of the configuration of spins in the Heisenberg model (Sec. ??). The implementation of the main steps in the DMQMC algorithm is described in more detail below.

2.7.1 Initial condition

The initial distribution of psips must be along the diagonal of the density matrix as this corresponds to a stochastic sampling of the identity matrix, which is the initial condition for the DMQMC algorithm. In order to create a psip on the diagonal, all spins in both bitstrings, that represent a psips position, are initially set down. Then spins are chosen at random and flipped up (in both bitstrings) until the configuration corresponds to the correct total spin value for the current simulation. As both bitstrings are equal, this must correspond to a diagonal element of the density matrix. The process is repeated until the correct number of initial pips have been spawned.

2.7.2 Spawning

The spawning is performed in a way that is very similar to in FCIQMC, but now a psip can attempt to spawn from both of its “spawning ends”, labelled A and B . One spawning end attempts to spawn onto a density matrix element in the same row, whilst the other end attempts to spawn on a density matrix element in the same column. In order to achieve this for the Heisenberg model, a single random excitationⁱ, is made to the bitstring that corresponds to the current spawning end.

The probability of spawning a psip on this new density matrix element, ρ_{ik} ,

ⁱFor the Heisenberg model the spawning corresponding Hamiltonian element is zero if the configurations differ by more than one excitation, so the spawning probability would be zero. An excitation is defined as flipping two adjacent and anti-parallel spins so that the total spin quantum number is conserved

from spawning end A is then given by

$$p_s(i, k|i, j) = \frac{1}{2} \frac{\Delta\beta |H_{jk}|}{P_{gen}(k|j)}, \quad (2.20)$$

where P_{gen} is the probability of generating the single excitation from $|j\rangle$ to $|k\rangle$. Here, the efficiency of the algorithm has been improved by only attempting to spawn to a single connected density matrix element. Hence, the re-weighting of the spawning probability by P_{gen} .

In a similar way, the probability of spawning a psip on a new density matrix element ρ_{kj} from spawning end B is given by

$$p_s(k, j|i, j) = \frac{1}{2} \frac{\Delta\beta |H_{ik}|}{P_{gen}(k|i)}. \quad (2.21)$$

In general, $\lfloor p_s \rfloor$ psips are always spawned and a further psip is spawned if $p_s - \lfloor p_s \rfloor > R$, where $0 \leq R \leq 1$ is a uniformly distributed pseudorandom numberⁱ.

2.7.3 Diagonal death/cloning step

The diagonal death/cloning step was performed in a similar way to the original FCIQMC code, but with a contribution from both spawning ends. For either of the spawning ends labelled by $|i\rangle$, the “probability” that the parent psip is killed by that spawning end is

$$p_d(i) = \frac{1}{2} \Delta\beta (H_{ii} - S). \quad (2.22)$$

In the current DMQMC implementation the probabilities of each spawning end are combined for simplicity. So for a density matrix element ρ_{ij} , the probability of a parent psip being killed is

$$p_d(i, j) = p_d(i) + p_d(j) = \left(\frac{H_{ii} + H_{jj}}{2} - S \right) \Delta\beta. \quad (2.23)$$

For $p_d \geq 0$, $\lfloor p_d \rfloor$ psips are always killed and a further psip is killed if $p_d - \lfloor p_d \rfloor > R$,

ⁱRandom numbers were generated using the `dSFMT` double-precision pseudorandom number generator[?].

with $0 \leq R \leq 1$ again being a uniformly distributed pseudorandom number. Moreover, in the rare event that $p_d < 0$, then $\lfloor |p_d| \rfloor$ psips are always cloned and a further psip is cloned if $|p_d| - \lfloor |p_d| \rfloor > R$.

2.7.4 Energy estimator

The calculation of the energy estimator is simplified if the Hamiltonian and the thermal density matrix are symmetric, which is the case for the Heisenberg Hamiltonian in Eq. ???. The expression for the estimator can be written

$$\langle H \rangle = \frac{\sum_{i,j} \rho_{ij} H_{ij}}{\sum_i \rho_{ii}}. \quad (2.24)$$

Now the denominator is trivial to calculate – it is the sum of the psip charges, q_{ii} , on the diagonal of the density matrix.

The numerator requires slightly more thought. As highlighted in Sec. ?? off-diagonal elements of the Heisenberg Hamiltonian, H_{ij} are only non-zero if the configurations $|i\rangle$ and $|j\rangle$ differ by a single excitation. In this case the density matrix element with total psip charge q_{ij} , contributes $-Jq_{ij}/2$ to the numerator of the energy estimator.

The two bitstrings that label the density matrix element can be compared, using the a bit-wise **XOR** operationⁱ. Half of the sum of the elements of the resultant bitstring is equal to the number of excitations. For example consider the two bitstrings 110100 and 011001. They differ by two excitations - to make them equal the two central spins and also the spins on both ends need to be flipped. In terms of the **XOR** operator this gives

$$110100 \text{ .XOR. } 011001 = 101101, \quad (2.25)$$

from which it is clear that the two bitstrings differ by two excitations.

The same bit-wise **XOR** operation can also be used to check whether an element is on the diagonal of the density matrix. In this case the contribution to the

ⁱIn more than one dimension, one must also test whether the two spins that are flipped between the two bitstrings are neighbours. This can be done by storing a list of neighbouring spins.

numerator of Eq. ?? is $q_{ii}H_{ii}$ where H_{ii} can be determined as described at the end of Sec. ??

2.7.5 Squared staggered magnetisation estimator

Staggered magnetisation is an order parameter of the antiferromagnetic Heisenberg model described in Sec. ?. An order parameter is a measure of the degree of order in a system and ranges between zero for total disorder and some maximum non-zero value for complete order. This measure usually changes during a phase transition. For a bipartite antiferromagnetic Heisenberg lattice that can be partitioned into two sublattices a and b such that neighbouring spins are on separate sublattices, the staggered magnetisation is given by[?]

$$m^\dagger = \sqrt{\langle (M^\dagger)^2 \rangle}, \quad (2.26)$$

$$M^\dagger = \frac{1}{N} \sum_{i=1}^N \epsilon_i S_i \quad (2.27)$$

$$\langle (M^\dagger)^2 \rangle = \frac{\sum_{i,j} (M^\dagger)_{ij}^2 \rho_{ji}}{\sum_i \rho_{ii}} \quad (2.28)$$

Where the co-efficient ϵ_i is $+1$ if the site i is on one sublattice and is -1 if site i is on the other sublattice. For the classical antiferromagnetic ground-state, or the Néel state, $m^\dagger = \frac{1}{2}$.ⁱ

It is easy to calculate the estimator for the squared staggered magnetisation if $(M^\dagger)^2$ is expressed in the form:

$$(M^\dagger)^2 = \frac{1}{N^2} \sum_{i,j} \epsilon_i \epsilon_j (S_i^x S_j^x + S_i^y S_j^y + S_i^z S_j^z). \quad (2.29)$$

Now for diagonal elements of $(M^\dagger)^2$, the x and y spin operator terms only contribute for $i = j$. If this is the case then both terms contribute $1/4N$ to the diagonal element. The contribution from the z spin operator term can be deduced from the total number of up spins, N_\uparrow , and the total number of up spins

ⁱAn interesting point is the difference in order for a quantum ground-state and the classical Néel state. When no external magnetic field is applied, the order of the quantum ground-state is zero[?]. This is as a result of intrinsic zero-point quantum spin fluctuations in the ground-state[?]

on sublattice a , N_{\uparrow}^a . Overall the diagonal elements are given by

$$(M^{\dagger})_{ii}^2 = \frac{(2N_{\uparrow}^{a,i} - N_{\uparrow})^2 + \frac{N}{4}}{N^2}, \quad (2.30)$$

where $N_{\uparrow}^{a,i}$ is the number of up spins on sublattice a for configuration $|i\rangle$. In the current implementation, this calculation of $N_{\uparrow}^{a,i}$ is made highly efficient by performing a bit-wise AND operation on the bitstring representation of $|i\rangle$ with a boolean mask for subsystem a . The sum of the elements of the resultant bitstring is equal to $N_{\uparrow}^{a,i}$.

For off-diagonal elements, $(M^{\dagger})_{ij}^2$, the contributions from the x and y spin operator terms are non-zero if $|i\rangle$ and $|j\rangle$ differ by one (not necessarily nearest-neighbour) excitation. If this is the case then $(M^{\dagger})_{ij}^2 = \pm 1/N^2$ depending upon whether the two spins that have been flipped are on the same lattice. Whether they are on the same lattice can be checked very quickly using the boolean mask for system a .

So in a similar way to the energy estimator, a density matrix element only contributes to the numerator of Eq. ?? if there are zero or one excitations between the two bitstrings that label the element. Then the contribution is given by $q_{ij}(M^{\dagger})_{ij}^2$, where $(M^{\dagger})_{ij}^2$ is found using the above prescription.

2.7.6 Heat capacity estimators

The fundamental equation of thermodynamics for a system of temperature T that is subject to a uniform magnetic field, h , in the z -direction can be written,

$$dE = dQ + dW = TdS - M_z dh, \quad (2.31)$$

where dE is the infinitesimal change in internal energy, dS is the infinitesimal change in entropy for a small flow of heat dQ , and $M_z dh$ is the work done on the system's total magnetic dipole moment \vec{M} by the magnetic field.

The heat capacity, C , can be defined by $dQ = CdT = TdS$. Assuming that the external magnetic field is constant then

$$dE = C_h dT, \quad (2.32)$$

where C_h is now the heat capacity for a constant magnetic field.

Expressing this in terms of β one obtains

$$C_h = -\beta^2 \frac{dE}{d\beta}. \quad (2.33)$$

At this point there is already a route for calculating the heat capacity in a constant magnetic field. DMQMC can easily sample the internal energy, so it is possible to simply differentiate this estimator in order to obtain the finite-temperature spectrum for C_h ,

$$\langle C_h \rangle = -\beta^2 \frac{d\langle H \rangle}{d\beta}. \quad (2.34)$$

In practice the data contains a lot of statistical noise, which makes it difficult to accurately approximate the gradient function of $\langle H \rangle$. However, a smoothing spline fit can be approximated for $\langle H \rangle$ in order to smooth out any statistical fluctuations (Appdx. B). From this the derivative in ?? can also be approximated.

Another, perhaps more elegant, method would be to sample C_h directly during the DMQMC simulation. If $\langle H \rangle$ is expanded into an orthonormal basis of energy eigenstates $\{|E_i\rangle\}$,

$$\langle H \rangle = \frac{\text{Tr}(\rho H)}{\text{Tr} \rho} = \frac{\sum_i \langle E_i | \rho H | E_i \rangle}{\sum_i \langle E_i | \rho | E_i \rangle} = \frac{\sum_i \langle E_i | e^{-\beta E_i} E_i | E_i \rangle}{\sum_i \langle E_i | e^{-\beta E_i} | E_i \rangle}, \quad (2.35)$$

it is simple to calculate a value of C_h using only the density matrix and the hamiltonian;

$$\begin{aligned} C_h &= \beta^2 \frac{d}{d\beta} \left[\frac{\sum_i \langle E_i | e^{-\beta E_i} E_i | E_i \rangle}{\sum_i \langle E_i | e^{-\beta E_i} | E_i \rangle} \right] \\ &= \beta^2 \frac{\text{Tr}(\rho H^2) \text{Tr}(\rho) - \text{Tr}(\rho H)^2}{\text{Tr}(\rho)^2} \\ &= \beta^2 (\langle H^2 \rangle - \langle H \rangle^2). \end{aligned} \quad (2.36)$$

Where the differentiation is performed using the quotient rule. This is in a form that can be sampled directly during the DMQMC simulation.

It is not immediately obvious whether direct stochastic sampling or a spline fit derivative is best for obtaining $\langle C_h \rangle$, so both will be tested in Sec. ??.

2.7.7 Other estimators

During the project estimators for spin-spin correlation and squared energy (required for the directly sampled heat capacity estimator) were also implemented. However, there is not enough room in this report to include the details or results for these estimators.

2.7.8 Statistical error calculations

A general DMQMC estimator can be written in the form,

$$O = \frac{x}{y}, \quad (2.37)$$

where x and y are separately averaged over all β -loops. As x and y are not necessarily uncorrelated, the standard error for the estimator was found by propagating the standard errors in the numerator and denominator using

$$\frac{\epsilon_O}{O} = \sqrt{\left(\frac{\epsilon_x}{x}\right)^2 + \left(\frac{\epsilon_y}{y}\right)^2 - 2\frac{\text{Cov}_{xy}}{xy}}, \quad (2.38)$$

where $\{\epsilon_i\}$ are the standard errors and Cov_{xy} is the covariance between x and y . This error analysis was performed in a post-processing step using a script written in `PYTHON` in order to avoid unnecessary calculations during the simulation.

Additionally, it is possible to obtain estimators for ground-state properties. One can average a finite-temperature estimator over consecutive β -steps once the simulation has settled into the ground-state. The Flyvbjerg and Peterson blocking analysis[?] can be used to take into account the correlations between consecutive β -steps when calculating the statistical error. During this project James Spencer's `PYTHON` implementation of this blocking analysis (originally applied to FCIQMC calculations) was employed.

2.8 Summary

In this chapter the DMQMC algorithm has been developed from an analogy with the imaginary-time Schrödinger equation and FCIQMC. In theory it should overcome the main two restrictions of FCIQMC, but may be restrained by the fact that the psips now live in a space that is the size of the Hilbert space of spin configurations squared. Additionally the implementation of the DMQMC algorithm was discussed, along with its application to the Heisenberg model.

Chapter 3

Testing and Results

In this chapter the implementation of DMQMC is tested on some small antiferromagnetic bipartite Heisenberg lattices. The benefits of testing on such systems is that they are sign-problem free[?] and the ground-state properties have been the subject of extensive study in the literature. All results in this section were simulated in the $M_S = 0$ total spin subspace with an inverse-temperature step of $\Delta\beta J = 0.1$ and in the absence of an external magnetic field ($h = 0$).

3.1 The 4×4 antiferromagnetic Heisenberg Lattice

DMQMC was first tested on the well-studied 4×4 antiferromagnetic Heisenberg lattice. The Hilbert space contains only 12,870 configurations in the $M_S = 0$ subspace, which allows calculation of the exact ground-state properties using direct methods such as FCI.

3.1.1 Energy

The DMQMC finite-temperature energy estimator for the 4×4 lattice is shown in Fig. ?? and Fig. ?. DMQMC mainly shows a good agreement with the exact FCI finite-temperature energy, which was obtained via a direct Lanczos diagonalisation[?]. However, there is a slight bias for $0.25 - 1 \beta J$. It was later later

discovered that this was a problem with accuracy and that the bias disappears if the inverse-temperature step is reduced to $\Delta\beta = 0.01$.




Figure 3.1: The DMQMC finite-temperature energy estimator (blue) and the exact FCI energy (red) for the 4×4 antiferromagnetic Heisenberg model. The DMQMC simulation ran with an initial population of 10^5 psips and for 10^3 β -loops.




Figure 3.2: The DMQMC finite-temperature energy estimator with statistical errors (blue) and the exact FCI energy (red) for the 4×4 antiferromagnetic Heisenberg model.

A blocking analysis over a single β -loop yields a ground-state energy estimate

of $E_0 = -0.7021(3)$, which is within error of the FCI ground-state energy[?] of $E_0 = -0.7018$.

However, the number of psips used was larger than the size of the Hilbert spaceⁱ. So whilst the DMQMC results are reasonably accurate, the finite-temperature properties can be calculated with far less effort using FCI. For example, the DMQMC simulation running on 64 cores (2.8 GHz) took approximately 3 hours to complete, whereas the FCI calculation took approximately 30 minutes on 2 cores (1.86 GHz). This could prove to be a problem with larger systems where even more psips are required.

3.1.2 Staggered magnetisation

The DMQMC finite-temperature squared staggered magnetisation estimator was also tested on the 4×4 antiferromagnetic Heisenberg model and the results are presented in Fig. ?? and Fig. ?. A blocking analysis over a single β -loop yields a ground-state squared staggered magnetisation estimate of $\langle (M^\dagger)^2 \rangle_0 = 0.2762(5)$, which agrees (within error) with the exact result of $\langle (M^\dagger)^2 \rangle_0 = 0.2765$ obtained by Dagotto and Moreo[?] using a modified Lanczos method.

ⁱBut considerably less than the 1.7×10^8 elements of the density matrix




Figure 3.3: The DMQMC finite-temperature squared staggered magnetisation estimator (blue) and the exact modified Lanczos ground-state squared staggered magnetisation (red) for the 4×4 antiferromagnetic Heisenberg model. The DMQMC simulation ran with an initial population of 10^5 psips and for 10^3 β -loops.




Figure 3.4: The DMQMC finite-temperature squared staggered magnetisation estimator with statistical errors (blue) and the exact modified Lanczos ground-state squared staggered magnetisation (red) for the 4×4 antiferromagnetic Heisenberg model.

3.1.3 Heat capacity

Two methods for calculating the heat capacity in a constant magnetic field were introduced in Sec. ???. One method was to calculate the derivative in Eq. ?? using a cubic spline fit of the finite-temperature energy estimator and the other was to directly sample the heat capacity during the DMQMC simulation. Both were tested on the 4×4 antiferromagnetic Heisenberg lattice with no magnetic field and are shown in Fig. ??.




Figure 3.5: The heat capacity estimator calculated by the spline derivative method (red) and by direct stochastic sampling (blue) for a 4×4 antiferromagnetic Heisenberg lattice.

It was found that the spline fit method worked best. The direct stochastic sampling of the heat capacity suffered from huge fluctuations as β grew large. This is due to the β^2 factor in Eq. ???. As β becomes large the fluctuations in $\langle H^2 \rangle - \langle H \rangle^2$ are amplified resulting in heavy statistical noise. This is not the case for the other estimators studied, where the factor is absent. The spline fit on the $\langle H \rangle$ estimator smooths out the statistical fluctuations and removes the dependence on the $\langle H^2 \rangle$ estimator, so performs better. Furthermore, the

spline derivative method tends towards the correct value in the large β limit. The heat capacity should be zero at absolute zero otherwise the third law of thermodynamics would be violated

3.2 The 6×6 antiferromagnetic Heisenberg Lattice

The Hilbert space of the 6×6 antiferromagnetic Heisenberg lattice contains a much larger 9.08×10^9 configurations in the $M_S = 0$ subspace. Although it is unfeasible to calculate exact FCI results for this system, it has been well-studied in the literature by methods such as GFQMC.

Fig. ?? shows the DMQMC finite-temperature energy estimator fluctuating wildly after only a short range in β . It was not possible to obtain the ground-state energy. This simulation took roughly 10.5 hours on 64 cores (2.8 Ghz) and although the number of psips was only a fraction of the Hilbert space, DMQMC cannot be regarded as a competitive QMC algorithm unless some major improvements are found.




Figure 3.6: The DMQMC finite-temperature energy estimator for the 6×6 antiferromagnetic Heisenberg model. The simulation ran with an initial population of 10^6 psips and for 10^3 β -loops.

3.3 Importance sampling

It was soon discovered that DMQMC was ineffective because the number of psips on the diagonal of the density matrix decays exponentially with inverse-temperature as shown in Fig. ?? . As psips spawn out from the diagonal and start to explore the complete space of density matrix elements, the chance of them spawning back onto the trace becomes very small. For example, psips that carry more than 1 excitation between their bitstring labels can no longer spawn directly back onto the diagonal, but have to do so in a number of steps. Meanwhile, the psips that were originally on the trace are killed off by the clone/death step of the DMQMC algorithm and so the population quickly diminishes.




Figure 3.7: Total number of psips on the trace as a function of β .

In order to improve the statistics on the trace at large β , spawning events onto the diagonal can be encouraged in a way that can be undone when calculating expectation value estimators. This can be achieved by increasing the probability of psips spawning onto the diagonal and assigning them a weight to be taken into account when calculating the estimators. More generally, to improve the

statistics across all excitations, psips can either be encouraged or discouraged to spawn from one excitation to another and then assigned the relevant weight.

One can think of a scalar function \mathcal{M} , which maps an element of $\boldsymbol{\rho}$ to an element of a new matrix $\boldsymbol{\rho}'$, which evolves with improved statistics across all excitations. The scalar function takes as its arguments, the configurations $|i\rangle$ and $|j\rangle$ that represent the position of the element in $\boldsymbol{\rho}$ and the matrix element ρ_{ij} itself. The function compares the two configurations and determines the number of excitations between them and hence scales the matrix element accordingly (equivalent to scaling the number of psips). So \mathcal{M} can be thought of in the following way,

$$\rho'_{ij} = \mathcal{M}(\rho_{ij}, |i\rangle, |j\rangle) = W(|i\rangle, |j\rangle)\rho_{ij}, \quad (3.1)$$

where W is a scalar function that determines the excitation between $|i\rangle$ and $|j\rangle$ and returns the appropriate scaling factor. It is important to note that $W(|i\rangle, |j\rangle) = W(|j\rangle, |i\rangle)$ as $\boldsymbol{\rho}$ is symmetric.

Now the inverse of \mathcal{M} is obvious,

$$\mathcal{M}^{-1}(|i\rangle, |j\rangle, \rho'_{ij}) = \frac{1}{W(|i\rangle, |j\rangle)}\rho'_{ij} = \rho_{ij} \quad (3.2)$$

To make notation simple define

$$\mathcal{M}(\rho_{ij}) \equiv \mathcal{M}(|i\rangle, |j\rangle, \rho_{ij}) \quad (3.3)$$

$$\mathcal{M}^{-1}(\rho'_{ij}) \equiv \mathcal{M}^{-1}(|i\rangle, |j\rangle, \rho'_{ij}) \quad (3.4)$$

$$W_{ij} \equiv W(|i\rangle, |j\rangle). \quad (3.5)$$

Under this transformation the estimator for the expectation value of an operator, O , can be written

$$\langle O \rangle = \frac{\sum_{i,j} \rho_{ij} O_{ji}}{\sum_i \rho_{ii}} = \frac{\sum_{i,j} \mathcal{M}^{-1}(\rho'_{ij}) O_{ji}}{\sum_i \mathcal{M}^{-1}(\rho'_{ii})} \quad (3.6)$$

$$= \frac{\sum_{i,j} \frac{1}{W_{ij}} \rho'_{ij} O_{ji}}{\sum_i \mathcal{M}^{-1}(\rho'_{ii})} \quad (3.7)$$

$$= W_{ii} \frac{\sum_{i,j} \rho'_{ij} \mathcal{M}^{-1}(O_{ji})}{\sum_i \rho'_{ii}}. \quad (3.8)$$

In the final step $\frac{1}{W_{ij}}O_{ji} = \mathcal{M}^{-1}(O_{ji})$ due to the symmetry of W and the fact that \mathbf{O} and $\boldsymbol{\rho}$ are expressed in the same basis. Moreover, W_{ii} has been taken out of the summation because for each i the excitation level is zero and hence W_{ii} is constant.

So under the importance sampling transformation (Eq. ??), expectation value estimators can be obtained by applying the inverse transformation to the matrix \mathbf{O} , then taking the expectation value with respect to $\boldsymbol{\rho}'$ and finally multiplying by the weight given to the zeroth excitation. If the weights W_{ij} are chosen carefully then the statistical error on the expectation value can be minimised.

3.3.1 New population dynamics

Given the importance sampling transformation (Eq. ??) it is possible to deduce a new set of rules for the population dynamics for the evolution of $\boldsymbol{\rho}'$, by transforming Eq. ??,

$$\begin{aligned} \mathcal{M}^{-1}(\rho'_{ij}(\beta + \Delta\beta)) = \\ \mathcal{M}^{-1}(\rho'_{ij}(\beta)) + \frac{1}{2} \sum_m (T_{im} \mathcal{M}^{-1}(\rho'_{mj}(\beta)) + \mathcal{M}^{-1}(\rho'_{im}(\beta)) T_{mj}) \Delta\beta. \end{aligned} \quad (3.9)$$

This simplifies to

$$\rho'_{ij}(\beta + \Delta\beta) = \rho'_{ij}(\beta) + \frac{1}{2} \sum_m \left(\frac{W_{ij}}{W_{mj}} T_{im} \rho'_{mj}(\beta) + \rho'_{im}(\beta) \frac{W_{ij}}{W_{im}} T_{mj} \right) \Delta\beta \quad (3.10)$$

Therefore the rates at which psips spawn has changed so that they spawn from

$$(i, m) \rightarrow (i, j) \quad \text{with probability} \quad \frac{W_{ij}}{2W_{im}} |T_{mj}| \quad (3.11)$$

and from

$$(m, j) \rightarrow (i, j) \quad \text{with probability} \quad \frac{W_{ij}}{2W_{mj}} |T_{im}|. \quad (3.12)$$

It is also important to note that when spawning in the opposite direction to in Eq. ?? and Eq. ?? then the inverse weights must be applied.

3.3.2 Choosing the weights

An optimal set of weights can be chosen by studying the number of psips on each excitation level during a standard DMQMC run. Fig. ?? shows an example how the population of psips on each excitation changes as a function of β . In the ground-state some excitation levels are poorly sampled because they are occupied by small numbers of psips.




Figure 3.8: The distribution of psips across the different excitation levels for a single β -loop of a DMQMC simulation of a 4×4 antiferromagnetic Heisenberg lattice.

In the ground-state the number of psips on each excitation remains approximately constant. At this point the weights from Sec. ?? can be determined so that the total population of psips is distributed uniformly over the different excitation levels. For instance,

$$W_n = \frac{N_p^{tot}}{N_{ex} N_p^n}, \quad (3.13)$$

where W_n is the weight for the n^{th} excitation, N_p^{tot} is the total number of psips

across all excitations, N_{ex} is the total number of excitations and N_p^n is the total number of psips on the n^{th} excitation.

When the importance sampling method was first tested, the weights were applied at the start of the simulation as shown in Fig ??.




Figure 3.9: The distribution of psips across the different excitation levels for a single β -loop of a DMQMC simulation of a 4×4 antiferromagnetic Heisenberg lattice, with importance sampling weights fully applied at $\beta = 0$.

It is evident that initially spawning events from the diagonal are restricted leading to an effective truncation of the Hilbert space as other excitations are suppressed. This meant that the higher excitations were being under-sampled for the initial evolution, which lead to biased results.

The weights were then introduced gradually, as shown in Fig. ??, so that in the ground-state ($\beta > \beta_{gs}$) they had reached the intended values. Each weight, W , was set to unity at $\beta = 0$ and for each step of $\Delta\beta$ they were scaled by $W^{\frac{\Delta\beta}{\beta_{gs}}}$ so that for $\beta > \beta_{gs}$ the weights were fully applied. Gradually introducing the weights in this way avoided a discontinuity in the spawning probabilities, whilst also reducing the effective truncation of the Hilbert space. In this way, the psip

populations in the pre-ground-state evolution is more similar to the case where no importance sampling is applied (Fig. ??).




Figure 3.10: The distribution of psips across the different excitation levels for a single β -loop of a DMQMC simulation of a 4×4 antiferromagnetic Heisenberg lattice, with the importance sampling weights introduced gradually.

3.3.3 Importance sampling implementation

This method of importance sampling is simple to implement. When attempting to spawn, as in Eq. ??, the excitation level of the current density matrix element can be found by performing a bit-wise **XOR** between the two bistrings that label it. The number of bits that are set in the output is equal to the number of excitations between the two bitstrings. This operation can also be used to determine the number of excitations between the two bitstrings labelling the density matrix element that a psip is spawning to. It is then easy to apply the correct weight to the spawning probability using Eq. ??.

When adding the contribution from a particular density matrix element to the estimator, the above weighting can be undone easily by calculating the excitation

level between the two labelling bit strings (using the bit-wise XOR operation) and then dividing by the corresponding weight.

3.4 The 6×6 antiferromagnetic Heisenberg lattice revisited

Now that an importance sampling method has been established, the 6×6 antiferromagnetic Heisenberg model can be re-investigated.

3.4.1 Energy

The importance-sampled DMQMC energy estimator is shown in Fig ?? and with errors in Fig ?. It shows a dramatic improvement in performance over standard DMQMC and with considerably less sampling. The energy now converges towards a reasonable value for the ground-state energy. For example, a blocking analysis of the DMQMC energy estimator yielded a ground-state energy of $E_0 = -0.6781(5)$, compared to Runge's GFQMC ground-state energy[?] of $E_0 = -0.678871(8)$. An interesting point to note is that in this case the initial number of psips occupy only 1.21×10^{-13} of the elements of the density matrix, so the performance improvement is huge.

3.4.2 Staggered magnetisation

The importance-sampled DMQMC squared staggered magnetisation estimator is shown in Fig ?. The squared staggered magnetisation estimator converges to the ground-state value of $\langle (M^\dagger)^2 \rangle_0 = 0.2098(2)$, which was found via a blocking analysis of a single β -loop. This agrees with the value of $\langle (M^\dagger)^2 \rangle_0 = 0.20986(8)$ that was calculated by Runge using GFQMC[?].




Figure 3.11: The DMQMC finite-temperature energy estimator with importance sampling (blue) and an accurate GFQMC estimate for the ground-state energy (red) for the 6×6 antiferromagnetic Heisenberg model. The simulation ran with an initial population of 1×10^7 psips and for 6 β -loops.




Figure 3.12: The importance-sampled DMQMC finite-temperature energy estimator with statistical errors (blue) and an accurate GFQMC estimate for the ground-state energy (red) for the 4×4 antiferromagnetic Heisenberg model.




Figure 3.13: The DMQMC finite-temperature squared staggered magnetisation estimator with importance sampling (blue) and an accurate GFQMC estimate for the ground-state squared staggered magnetisation (red) for the 6×6 antiferromagnetic Heisenberg. The simulation ran with an initial population of 10^7 psips and for 6 β -loops.




Figure 3.14: The importance-sampled DMQMC finite-temperature squared staggered magnetisation estimator with statistical errors (blue) and an accurate GFQMC estimate for the ground-state square staggered magnetisation (red) for the 6×6 antiferromagnetic Heisenberg model.

3.5 The $4 \times 4 \times 4$ antiferromagnetic Heisenberg lattice

The importance-sampled DMQMC algorithm was also tested on the $4 \times 4 \times 4$ antiferromagnetic Heisenberg lattice and the results are shown in Fig. ?? and Fig. ?. It was difficult to test the accuracy of the DMQMC method in three dimensions, as the three dimensional antiferromagnetic Heisenberg model features less prominently in the literature. Nevertheless, the DMQMC ground-state energy estimate was compared against an accurate ground-state energy estimate obtained via FCIQMC[?].

Performing a blocking analysis, on a single β -loop DMQMC simulation gives a ground-state energy of $E_0 = -0.912(6)$, whereas a single run of the FCIQMC algorithm[?] gives a ground-state energy of $E_0 = -0.91509(5)$. The FCIQMC value was obtained using an initial population of 10^6 psips. DMQMC seems a good compromise between accuracy and the ability to obtain finite-temperature estimates.

The Hilbert space for a $4 \times 4 \times 4$ lattice contains roughly 1.83×10^{18} spin configurations and so in this case DMQMC was simulating a density matrix with approximately 1.83×10^{36} elements using only 5×10^5 psips. The statistical errors in Fig. ?? are surprisingly small considering the number of psips used.




Figure 3.15: The DMQMC finite-temperature energy estimator with importance sampling (blue) and an accurate FCIQMC estimate for the ground-state energy (red) for the $4 \times 4 \times 4$ antiferromagnetic Heisenberg model. The simulation ran with an initial population of 5×10^6 psips and for 43 β -loops.




Figure 3.16: The importance-sampled DMQMC finite-temperature energy estimator with statistical errors (blue) and an accurate FCIQMC estimate for the ground-state energy (red) for the $4 \times 4 \times 4$ antiferromagnetic Heisenberg model.

3.6 Summary

The DMQMC implementation was tested on some simple antiferromagnetic Heisenberg models in both two and three dimensions. With the standard algorithm, developed in Ch. ??, it was found that the initial population of psips was similar to the size of the Hilbert space of configurations and failed for larger systems. It appeared as though the method might not be efficient enough to compete as a QMC method.

However, the reason for this poor performance was soon realised and a suitable importance sampling adaptation to the algorithm was invented, giving dramatically improved results. This new method allowed for accurate finite-temperature calculations of energy and staggered magnetisation for the 6×6 Heisenberg lattice, for which the conventional DMQMC method had failed. Importance sampling also enabled accurate simulation of the $4 \times 4 \times 4$ antiferromagnetic Heisenberg model. In most cases the DMQMC ground-state estimators (obtained via a blocking analysis) were equal to the results in the literature, within error.

Additionally, both versions of the heat capacity estimator were tested on the 4×4 antiferromagnetic Heisenberg lattice. It was found that directly sampling the heat capacity from the DMQMC simulation gave poor results, even with importance sampling. This was because the statistical fluctuations were amplified for large β . This proved the point that, although in DMQMC it is theoretically possible to calculate the expectation value of any quantum mechanical operator, in practice it can prove to be difficult to obtain accurate results. The spline derivative method did provide improved results, but this method requires post-processing of the data and makes it difficult to determine errors.

Chapter 4

Application to Quantum Information

Entanglement is the fundamental inseparability of quantum systems - seen as correlations that emerge in quantum theory and that cannot be described by classical physics. It has been described as “the strangest and least understood facet of quantum theory”[?] and this, in itself, provides a good motivation for the study of quantum entanglement. Furthermore, some future technologies such as quantum computers, quantum teleportation and ion clocks[?] rely on the unique properties of entanglement. A method that can calculate measures of entanglement between qubits or systems of qubits could prove useful in the development of such technologies.

Currently, there are no known scalable simulation methods[?] that are capable of calculating entanglement measures. Entanglement simulations rely on non-scalable algorithms such as density matrix renormalisation group (DMRG) or Lanczos. Other QMC methods do not provide access to the reduced density matrix. However, with DMQMC and the stochastic sampling of the full density matrix, it is possible to obtain the reduced density matrix.

This chapter begins by introducing the two main measures of entanglement (concurrence and Von-Neumann entropy) in quantum information theory. It moves on to discuss how estimators for such methods can be incorporated into the DMQMC algorithm. Finally, the estimator for concurrence is tested against

some simple 1D antiferromagnetic Heisenberg rings in the presence and absence of an applied magnetic field.

4.1 Reduced Density Matrix

A major obstacle that prevents QMC simulations of entanglement measures in many-electron quantum systems is the difficulty of obtaining the reduced density matrix. In DMQMC, it is possible to obtain a reduced density matrix for a small subsystem (small enough so that it can be stored and manipulated in computer memory).

Consider two quantum systems A and B that are coupled to form a system C (Fig. ??) such that

$$\mathcal{H}_C = \mathcal{H}_A \otimes \mathcal{H}_B, \quad (4.1)$$

where \mathcal{H}_A , \mathcal{H}_B and \mathcal{H}_C denote the Hilbert spaces of systems A , B and C .




Figure 4.1: Two quantum systems, A and B coupled such that the quantum state of the complete system lives in a Hilbert space that is the tensor product of the Hilbert spaces of A and B .

In the Schmidt basis[?] a quantum state, $|\psi\rangle$, of system C can be expressed in terms of the basis vectors of systems A and B ,

$$|\psi\rangle = \sum_n \alpha_n |a_n\rangle \otimes |b_n\rangle. \quad (4.2)$$

With $|a_n\rangle \in \mathcal{H}_A$ and $|b_n\rangle \in \mathcal{H}_B$. Now consider taking a measurement, O_A , on

system A without touching B , the expected outcome of this measurement is then,

$$\langle \psi | O_A \otimes \mathbb{1}_B | \psi \rangle = \langle \psi | \sum_n \alpha_n O_A | a_n \rangle \otimes | b_n \rangle \quad (4.3)$$

$$= \sum_{n,m} \alpha_m^* \alpha_n \langle a_m | O_A | a_n \rangle \otimes \langle b_m | b_n \rangle \quad (4.4)$$

$$= \sum_n |\alpha|^2 \langle a_n | O_A | a_n \rangle \quad (4.5)$$

$$= \text{Tr}_A(\rho_A O_A), \quad (4.6)$$

where $\mathbb{1}_B$ is the identity operator acting on system B , ρ_A is the reduced density operatorⁱ for system A and Tr_A is the partial trace over the basis of system B . The reduced density operator, ρ_A may be obtained by “tracing out” system B from the projector onto state $|\psi\rangle \in \mathcal{H}_C$ i.e.

$$\rho_C = |\psi\rangle \langle \psi| = \sum_{n,m} \alpha_n \alpha_m^* |a_n\rangle \langle a_m| \otimes |b_n\rangle \langle b_m| \quad (4.7)$$

$$\rho_A = \text{Tr}_B(\rho_C) = \sum_{n,m} \alpha_n \alpha_m^* |a_n\rangle \langle a_m| \otimes \text{Tr}(|b_n\rangle \langle b_m|) \quad (4.8)$$

$$= \sum_n |\alpha|^2 |a_n\rangle \langle a_n|. \quad (4.9)$$

So in DMQMC, it is also possible to stochastically sample the reduced density matrix for a subsystem of the total system under study. From this it is possible to calculate expectation values for measurement on that subsystem. Moreover, it allows for some interesting measures of quantum entanglement between the subsystems A and B .

4.2 Von-Neumann Entropy

The Von-Neumann Entropy is an extension of the Shannon entropy, from classical information theory into quantum information theory. In classical information

ⁱit satisfies the three properties in Sec. ??

theory, the Shannon entropy for a probability distribution q_1, \dots, q_n is

$$S_C(q_1, \dots, q_n) = \sum_i q_i \log_2 1/q_i. \quad (4.10)$$

Similarly for a quantum state living in a finite D -dimensional Hilbert space, the Von-Neumann entropy is

$$S_Q(\boldsymbol{\rho}) = -\text{Tr}(\boldsymbol{\rho} \log_2 \boldsymbol{\rho}) = \sum_i p_i \log_2 1/p_i, \quad (4.11)$$

where p_i are the eigenvalues of $\boldsymbol{\rho}$. It describes the purity of the system. If one considers the density matrix in diagonal form, the eigenvalues p_i represent the probability of the system being in some state $|i\rangle$ of the diagonal basis. Hence, for a pure state one eigenvalue is unity, whilst all the others are zero, giving a zero Von-Neumann entropy. However, for a maximally entangled mixed state, with equally likely probabilities of being in any of the states $|i\rangle$ in the diagonal basis, the Von-Neumann entropy is $\log_2 D$.

Therefore, if the Von-Neumann entropy is calculated for the reduced density matrix for system A , in Fig. ??, it measures the amount of quantum entanglement between systems A and B . For example, if the Von-Neumann entropy is zero, then A is a pure state and so there is no entanglement between A and B , assuming that the overall system C is in a pure state. On the other hand if the Von-Neumann entropy is $\log_2 D$, then subsystem A is in a fully mixed state and A and B must be maximally entangled.

In order to compute the Von-Neumann entropy (Eq. ??), the subsystem A must be sufficiently small so that the reduced density matrix can be diagonalised directly. The Von-Neumann entropy has not been properly tested due to time restrictions.

4.3 Concurrence and Entanglement of Formation

In 2001, Hill and Wootters[?] introduced a new measure of entanglement for a two-qubit mixed state. For the case where the total system is mixed, the Von-Neumann entropy is no longer a good measure of quantum entanglement, because both subsystems A and B can have non-zero entropy even if there is no entanglement. The measure, known as the “entanglement of formation” is a generalisation of Von-Neumann entropy. The entanglement of formation of a bipartite quantum system can be defined as “the minimum number of singlets needed to create an ensemble of pure states that represents a mixed state”[?]. In this case a “singlet” refers to a state with zero total spin.

Another definition of the Von-Neumann entropy, $E(\Phi)$, is the number of singlet pairs, required to required to create a copy of a pure state $|\Phi\rangle$ [?]. The number of singlets required to create a copy of some mixed state that is expressed as an ensemble of pure states, $\rho = \sum_{j=1}^N p_j |\Phi_j\rangle \langle \Phi_j|$, can be written

$$N_s = \sum_{j=1}^N p_j E(\Phi_j). \quad (4.12)$$

This number depends on the particular ensemble so the minimum over all pure-state decompositions is taken in order to pick out the irreducible entanglement of the mixed system

$$E_f(\rho) = \min \left(\sum_j^N p_j E(\Phi_j) \right). \quad (4.13)$$

For the special case of a two-qubit mixed state the entanglement of formation can be calculated from a quantity known as “concurrence”. For a reduced density matrix, ρ_A , where in this case A refers to a subsystem of two qubits, the concurrence is defined as

$$\mathcal{C}(\rho_A) \equiv \max(0, \gamma_1 - \gamma_2 - \gamma_3 - \gamma_4), \quad (4.14)$$

in which $\gamma_1 > \gamma_2 > \gamma_3 > \gamma_4$ are the eigenvalues of the matrix

$$\mathbf{R} = \sqrt{\sqrt{\boldsymbol{\rho}_A} \tilde{\boldsymbol{\rho}}_A \sqrt{\boldsymbol{\rho}_A}} \quad \text{with} \quad \tilde{\boldsymbol{\rho}}_A = (\sigma_y \otimes \sigma_y) \boldsymbol{\rho}_A^* (\sigma_y \otimes \sigma_y), \quad (4.15)$$

where σ_y is the Pauli spin matrix for the y -direction.

The concurrence, \mathcal{C} , ranges from zero to one and is monotonically related to the entanglement of formation[?]. In this way the concurrence can also be regarded as a measure of entanglement. The entanglement of formation for two qubits[?] is given by,

$$\mathcal{E}(\mathcal{C}) = h\left(\frac{1 + \sqrt{1 - \mathcal{C}^2}}{2}\right) \quad \text{with} \quad h(x) = -x \log_2 x - (1 - x) \log_2 (1 - x). \quad (4.16)$$

The concurrence is easy to compute for the Heisenberg model. If it is assumed that the Hamiltonian is real (as is the case in the Heisenberg model) and therefore that the reduced density matrix is also real, the problem is reduced to finding the moduli of the eigenvalues of

$$\tilde{\mathbf{R}} = \boldsymbol{\rho}_A (\sigma_y \otimes \sigma_y). \quad (4.17)$$

As $\tilde{\mathbf{R}}$ is only a 4×4 matrix it is trivial to compute the concurrence.

4.4 Implementation of the ground-state entanglement estimators

Whilst it is possible to calculate finite-temperature estimators for both the Von-Neumann entropy and Concurrence, the current implementation can only handle such estimators for the ground-state. This is because the current implementation is based on previous FCIQMC code which only allows simulations for a single M_S subspace. In FCIQMC the ground-state always lies within a known M_S . For example, in a simulation for the antiferromagnetic Heisenberg model with zero applied magnetic field the ground-state is within the $M_S = 0$ subspace. Unfortunately, there was not enough time during this MSci project to adapt the code so that simulations could run in all subspaces simultaneously, which would

allow for finite-temperature estimators of entanglement.

In the current implementation the reduced density matrix, for a specified subsystem of spins, is averaged over all $\beta > \beta_{gs}$, where β_{gs} is the critical β beyond which the simulation enters the ground-state. This is to ensure that the reduced density matrix has the highest possible accuracy before computing the entanglement estimators.

In order to obtain the reduced density matrix at a particular $\beta > \beta_{gs}$ the subsystem B is traced out as in Eq. ???. This was achieved by looping over all elements of the total density matrix with non-zero psip populations and determining whether a particular element contributed towards the reduced density matrix. A simple test (see Fig. ??) was devised to determine whether this was the case:

1. Define a boolean mask for subsystem B - this is created at the start of the simulation and has bits set to 1 where a spin is to be included in subsystem B .
2. Perform a bit-wise **AND** operation on each of the two bit string ends that label the current density matrix element, with the boolean mask for subsystem B . This sets all bits that represent a spin in subsystem A to 0 and leave those that represent a spin in B unchanged.
3. If the two resultant bit strings are equal then this means that it lies on the diagonal of subsystem B and hence contributes to the reduced density matrix for subsystem A . In this case, the psip charge on this element is added to the corresponding element of the reduced density matrix. If they are not equal then nothing is added to the reduced density matrix.




Figure 4.2: Illustration of how a given density matrix element ρ_{ij} is tested to see whether it contributes towards ρ_A for the case that A is 3 neighbouring spins on a chain of length $N = 8$

After obtaining the average reduced density matrix, the Von-Neumann entropy is obtained by directly diagonalising the reduced density matrixⁱ. For the concurrence, the matrix \tilde{R} (Eq. ??) is diagonalisedⁱⁱ. This whole process is repeated over multiple β -loops to build up statistics and to enable the calculation of errors as in Sec. ??.

4.5 Nearest-neighbour concurrence in Heisenberg rings

The ground-state concurrence for neighbouring spins (qubits), illustrated in Fig. ??, on a bipartite antiferromagnetic Heisenberg ring with N qubits was first calculated by O'Connor and Wothers in 2001[?]. They found that the ground-state concurrence, \mathcal{C}_{gs} , is related to the ground-state energy, E_0 , in the following way

$$\mathcal{C}_{gs} = -\frac{1}{2}(E_0/N + 1). \quad (4.18)$$

ⁱThis is achieved using the LAPACK `ssyev` and `dsyev` routines for diagonalising symmetric matrices. These routines use the QR algorithm.

ⁱⁱThis is achieved using the LAPACK `sgeev` and `dgeev` routines for diagonalising nonsymmetric matrices. Both routines first use Hessenberg reduction followed by the QR algorithm.




Figure 4.3: Illustration of the subsystem of two qubits, A , on an $N = 8$ antiferromagnetic Heisenberg chain. The system can be regarded as a ring due to the periodic boundary conditions adopted.

To test whether the DMQMC ground-state concurrence estimator was implemented correctly, several Heisenberg rings of differing size were simulated. The results in Tab. 4.1, show that the DMQMC ground-state concurrence estimator was correct within error for bipartite rings up to $N = 10$. Moreover, an $N = 36$ Heisenberg ring was simulated with importance sampling and the ground-state concurrence of $\mathcal{C}_{gs} = 0.3873(8)$ was obtained. This agrees well with the value of $\mathcal{C}_{gs} = 0.38748(4)$ obtained from FCIQMC[?] and Eq. ??.

N	2	4	6	8	10	...	36	∞
[?] \mathcal{C}_{gs}	1.000	0.500	0.434	0.412	0.403	...	–	0.386
DMQMC	1.0006(6)	0.5005(4)	0.4342(5)	0.4129(5)	0.4031(4)	...	0.3873(8)	–

Table 4.1: Comparison of the DMQMC ground-state concurrence calculation with the exact values for neighbouring qubits on a simple one dimensional bipartite antiferromagnetic Heisenberg ring of N qubits.

4.6 Concurrence in a uniform magnetic field

Applying a magnetic field to a system of spins changes the structure of the ground-state spin configuration and hence the ground-state entanglement properties of the system. In quantum information it is interesting to look at how the entanglement between qubits changes with a uniformly applied magnetic field. Since

the magnetic field can be controlled it is possible to also control the amount of entanglement between certain qubits.

As an example a uniform magnetic field was applied to a Heisenberg antiferromagnetic ring with $N = 8$, the effect of which is shown in Fig. ???. Within a given M_S subspace, changing the magnetic field shifts the diagonal elements of the Hamiltonian in Eq. ??. Therefore, although the energy of the ground-state is changing, the ground-state wave function does not. Hence, for a given M_S the concurrence between two neighbouring qubits does not change.




Figure 4.4: The ground-state concurrence, C_{gs} , between two neighbouring qubits on an $N = 8$ antiferromagnetic Heisenberg ring as a function of magnetic field strength in the (z -direction), h .

However, when the lowest energy of a particular M_S subspace is shifted so that it is higher than that of the next M_S , the ground-state wave function changes to one that belongs to the new M_S subspace. As the form of the wave function changes so do the entanglement properties of the system. This explains, why in Fig. ?? there are regions of constant entanglement.

As the magnetic field increases, the amount of entanglement decreases as more

spins are forced into alignment and the orientation of neighbouring spins becomes less correlated. Beyond a critical value of the magnetic field, $h/J = 0.125$, the concurrence drops to zero. This is the point where all spins become aligned and quantum correlations between neighbouring qubits disappears. It can be regarded as a quantum phase transition at $T = 0$ due to quantum fluctuations, rather than thermal ones[?].

4.7 Concurrence with increasing separation

Entanglement length[?], l_E , is another important property of many-body quantum systems that is relevant to quantum information theorists and experimentalists. It describes how quickly the entanglement between two qubits falls to zero as the separation between them is increased.

In the Heisenberg model it is assumed that spin-spin interactions are short-range and between nearest neighbours. It is therefore difficult to find quantum correlations between spins that are not nearest neighbours. However, there is another critical value of the magnetic field, h_{sym} , for which an antiferromagnetic Heisenberg system is in the symmetric state

$$\psi_{sym} = \frac{1}{\sqrt{N}}(|100 \dots 0\rangle + |010 \dots 0\rangle + |001 \dots 0\rangle + \dots + |000 \dots 1\rangle). \quad (4.19)$$

In this state it is evident that there is non-zero entanglement between any two qubits. This makes it possible to study how the entanglement between qubits changes as the distance between them is varied. Another benefit of this is that at this critical magnetic field the number of possible spin configurations is small and DMQMC can simulate the system very accurately.

Fig. ?? shows how the concurrence drops off as the distance between the two qubits increases, for an antiferromagnetic Heisenberg ring with $N = 36$. Ignoring the effects near $d = 0$ from the periodic boundary conditions, it can be seen that the concurrence approaches $\mathcal{C}_{gs} \propto e^{-d}$. Hence the entanglement length approaches $l_E \approx 1$. This makes sense as in the Heisenberg model it is approximated that the J -coupling occurs between nearest neighbours only and entanglement between non-nearest neighbours is due to non-direct interactions.




Figure 4.5: Ground-state concurrence with statistical errors (top) and logarithm of ground-state concurrence (bottom) for the $N = 36$ antiferromagnetic Heisenberg ring, with $\hbar/J = 1.37$, as a function of separation between the two qubits. For example, $d = 0$ is nearest neighbours, $d = 1$ is next-nearest neighbours etc. It was difficult to obtain good statistics for $|d| > 10$ as the concurrence was so small.

4.8 Summary

In this chapter two quantum entanglement measures, the Von-Neumann entropy and the concurrence, were introduced. And the integration of these entanglement estimators into the DMQMC algorithm was discussed. The ground-state concurrence estimator was tested on one dimensional antiferromagnetic Heisenberg rings and the results were found to agree very well with the analytic values.

Next the effects of a magnetic field on the entanglement between two nearest-neighbour qubits on an $N = 8$ antiferromagnetic Heisenberg ring were studied. It was found that there were quantum sharp transitions between different levels of entanglement due to quantum phase transitions. Finally, DMQMC was used to study how the entanglement between two qubits falls off as a function of the separation between them. It was found that there was a unit entanglement length, which agrees with the assumption in the Heisenberg model that interactions are only between nearest-neighbours.

Chapter 5

Conclusion

This project sought to overcome the two main restrictions of FCIQMC by stochastically sampling the many-electron thermal density matrix rather than the many-electron wave function. Although sampling the density matrix increases the computational effort because the psips now have to live, die, clone and annihilate in a space that has squared in size, it in theory allows for the calculations of any quantum mechanical observable at finite-temperatures.

The algorithm was developed by studying an analogy between the evolution of the many-electron wave function in imaginary-time and the many-electron thermal density matrix as a function of inverse-temperature. A set of population dynamics, similar to those in FCIQMC were chosen so that they simulate a first-order finite difference approximation that relate the thermal density matrix at one inverse-temperature to the thermal density matrix at the step before.

The DMQMC algorithm was implemented in `FORTRAN90` and built upon a current highly parallel and optimised implementation of FCIQMC. It was decided to test the implementation on some well-studied two and three dimensional antiferromagnetic bipartite Heisenberg lattices because they remove the complication of the fermion sign-problem which plagues some many-electron models. At first, it was found that the DMQMC method was far less efficient than FCI for the 4×4 antiferromagnetic Heisenberg lattice and actually failed for the 6×6 antiferromagnetic Heisenberg lattice.

It was found that the reason for this failure was that the number of psips on the trace decayed exponentially and quickly became undersampled. This was

remedied by the introduction of an importance sampling method that looked to increase the sampling of rare spawning events. Under this importance sampling transformation it was found that the number of initial psips could be significantly reduced and that DMQMC could now be used to simulate larger systems for which FCI is unfeasible.

Futhermore, larger systems such as the $4 \times 4 \times 4$ antiferromagnetic Heisenberg lattice could be simulated with a number of psips that was only a minute fraction of the total number of elements in the thermal density matrix. As expected, DMQMC was found to be less accurate than FCIQMC for calculating ground-state properties with a similar number of psips, but it has the benefit that all of data from the simulation was useful in determining finite-temperature properties. In FCIQMC relevant data can only really be retrieved once the system is in the ground-state in the limit of large imaginary-time.

In theory it is simple to calculate the expectation value of any quantum mechanical observable in DMQMC. However, in practice it was difficult to directly sample the estimator for the heat capacity. This was due to a factor of β^2 in the estimator equation that leads to the amplification of statistical errors as β becomes large. A solution was to calculate the derivative of the energy estimator using a cubic smoothing spline fit. However, this rather inelegant method leads to difficulties in obtaining statistical errors.

The ease of obtaining the reduced density matrix from stochastic sampling of the full thermal density matrix opened up the possibility of calculating entanglement measures and an avenue into quantum information theory. Current QMC methods do not provide access to the reduced density matrix and other methods are non-scalable or can only simulate one dimensional systems (DMRG). So there is need for methods such as DMQMC in quantum information. Two entanglement measures, Von-Neumann entropy and concurrence, were implemented of which the ground-state concurrence estimator was tested on one dimensional antiferromagnetic Heisenberg rings for which analytic values exist.

Some work on the DMQMC implementation is still needed. For example, the current implementation can only simulate systems within a fixed M_S subspace. One needs to find a way of either simulating across all M_S subspaces simultaneously or combining the results from all the different M_S simulations. At this

point, it would be possible to obtain the full finite-temperature picture across all M_S values and it would allow for finite-temperature entanglement calculations. Additionally, it seems that the FCIQMC shift update algorithm introduces biasesⁱ and this problem must be solved.

In conclusion, a new QMC method has been developed. It shares some of the important features of FCIQMC, but almost fully overcomes two of FCIQMC's main restrictions at the expense of efficiency. Some more work is required, but there are signs that it could be an important QMC method for finite-temperature calculations and could be used to tackle some poorly understood aspects of quantum information theory.

Total words: 9997 (via TeXcount 2.3)

ⁱThere was not enough space in this report to discuss this in full.

References

- . G. H. Booth, A. J. W. Thom, and A. Alavi, [J. Chem. Phys.](#) **131**, 054106 (2009).
- . A. Szabados, P. Jeszenszki, and P. R. Surján, [J. Chem. Phys.](#) , 1 (2011).
- . P. R. C. Kent, *Techniques and Applications of Quantum Monte Carlo*, [Ph.D. thesis](#), University of Cambridge (1999).
- . W. M. C. Foulkes, L. Mitas, R. J. Needs, and G. Rajagopal, [Rev. Mod. Phys.](#) **71**, 33 (2001).
- . M. Bajdich and L. Mitas, [Act. Phys. Slov.](#) **59**, 81 (2009), [arXiv:arXiv:1008.2369v1](#) .
- . J. S. Spencer, N. S. Blunt, and W. M. C. Foulkes, [J. Chem. Phys.](#) **136**, 054110 (2012).
- . C. J. Isham, *Lectures on Quantum Theory. Mathematical and Structural Foundations*, 1st ed. (Imperial College Press, 1995).
- . M. Foulkes and J. Spencer, Private communication (2011).
- . J. B. Anderson, [Physics](#) **1499** (1975), 10.1063/1.431514.
- . C. Contaldi, *Computational Physics Lecture Notes* (Imperial College London, 2011).
- . C. J. Umrigar, M. P. Nightingale, and K. J. Runge, [J. Chem. Phys.](#) **99**, 2865 (1993).
- . W. Heisenberg, *Zeitschrift fur Physik* **38** (1926).
- . P. Dirac, [Proceedings of the Royal Society of London.](#) **112**, 661 (1926).
- . M. Karbach and G. Muller, “Introduction to the Bethe ansatz I,” (1998), [arXiv:9809162 \[cond-mat\]](#) .
- . M. Saito and M. Matsumoto, in [Monte Carlo and Quasi-Monte Carlo Methods 2008](#), edited by P. L’ Ecuyer and A. B. Owen (Springer Berlin Heidelberg, Berlin, Heidelberg, 2009) pp. 589–602.
- . K. J. Runge, [Phys. Rev. B](#) **45**, 7229 (1992).
- . J. van den Brink, in *Lecture Notes on Theory of Condensed Matter* (University of Leiden, 2012) pp. 5–28.

REFERENCES

- . C. L. Henley, in *Lecture notes* (Cornell University, 2007) pp. 15–23.
- . H. Flyvbjerg and H. G. Petersen, *J. Chem. Phys.* **91**, 461 (1989).
- . Obtained using J. Spencer’s FCI code (2011).
- . E. Dagotto and A. Moreo, *Phys. Rev. B* **38**, 5087 (1988).
- . K. J. Runge, *Phys. Rev. B* **45**, 12292 (1992).
- . Obtained using J. Spencer’s FCIQMC code (2011).
- . T. Rudolph, *Quantum Information Lecture Notes* (Imperial College London, 2011).
- . M. B. Hastings, “Measuring Renyi Entanglement Entropy with Quantum Monte Carlo,” (2010), [arXiv:arXiv:1001.2335v2](#) .
- . A. Ekert and P. L. Knight, *Am. J. Phys.* **63**, 415 (1994).
- . S. Hill and W. Wootters, *Phys. Rev. Lett.* **78**, 5022 (1997).
- . W. Wootters, *Quant. Inf. Comp.* **1**, 27 (2001).
- . K. O’Connor and W. Wootters, *Phys. Rev. A* **63**, 1 (2001).
- . T. Rudolph, Private communication (2012).
- . M. Arnesen, S. Bose, and V. Vedral, *Phys. Rev. Lett.* **87**, 4 (2001).
- . M. Born and J. R. Oppenheimer, *Ann. Physik* **84**, 457 (1927).
- . A. Szabo and N. S. Ostlund, *Modern Quantum Chemistry: Introduction to Advanced Electronic Structure Theory* (Dover Publications, 1996).
- . B. H. Bransden and C. J. Joachain, *Quantum Mechanics*, 2nd ed. (Prentice Hall, 2000).
- . A. J. James, *PhD Thesis*, [Ph.D. thesis](#), Imperial College London (1996).
- . A. M. Mazzone and V. Morandi, *Comp. Mat. Sci.* **38**, 231 (2006).
- . S. Sorella, S. Baroni, R. Car, and M. Parrinello, *Europhys. Lett.* **8**, 663 (1989).
- . G. H. Booth and A. Alavi, *J. Chem. Phys.* **132**, 174104 (2010).
- . G. H. Booth, D. Cleland, A. J. W. Thom, and A. Alavi, *J. Chem. Phys.* **135**, 084104 (2011).
- . C. H. Reinsch, *Num. Math.* **83**.

Appendix A: Basic electronic structure and FCIQMC

This appendix briefly reviews some of the fundamentals in electronic structure theory and projector QMC methods before moving onto a brief outline of the recent FCIQMC method.

The many-electron Schrödinger Equation

For a many-electron problem with N electrons and M nuclei the Hamiltonian, in Hartree atomic units, is

$$\begin{aligned}
 H = & - \frac{1}{2} \sum_{i=1}^N \nabla_{\mathbf{r}_i}^2 - \sum_{i=1}^M \frac{1}{2M_A} \nabla_{\mathbf{r}_A}^2 - \sum_{i=1}^N \sum_{A=1}^M \frac{Z_A}{|\mathbf{r}_i - \mathbf{d}_A|} \\
 & + \frac{1}{2} \sum_{i=1}^N \sum_{j \neq i}^N \frac{1}{|\mathbf{r}_i - \mathbf{r}_j|} + \frac{1}{2} \sum_{A=1}^M \sum_{B \neq A}^M \frac{Z_A Z_B}{|\mathbf{d}_A - \mathbf{d}_B|}, \quad (1)
 \end{aligned}$$

where M_A is the ratio of the mass of the nucleus A to the mass of an electron; Z_A is the atomic number of nucleus A; \mathbf{r}_i is the position vector of the i th electron; \mathbf{d}_A is the position vector of the A th nuclei. The first term in Eq. ?? is the kinetic energy operator of the electrons; the second term is the kinetic energy operator of the nuclei; the third term represents the electron-nuclei coulomb attraction; the

fourth and fifth terms represent the electron-electron and nuclei-nuclei coulomb repulsions respectively.

The Born-Oppenheimer approximation[?] assumes that the nuclei are stationary, due to the fact they are much heavier than the electrons. Hence, the second term in Eq. ?? can be neglected and the nuclei-nuclei repulsion term is assumed to be constant. So the time-independent many-electron Schrödinger equation under the Born-Oppenheimer approximation becomes

$$\left[-\frac{1}{2} \sum_{i=1}^N \nabla_{\mathbf{r}_i}^2 - \sum_{i=1}^N \sum_{A=1}^M \frac{Z_A}{|\mathbf{r}_i - \mathbf{d}_A|} + \frac{1}{2} \sum_{i=1}^N \sum_{j \neq i}^N \frac{1}{|\mathbf{r}_i - \mathbf{r}_j|} \right] \psi(\mathbf{R}) = E \psi(\mathbf{R}). \quad (2)$$

The solution, $\psi(\mathbf{R})$, of Eq. ?? is the many-body wave function for N electrons and is a function of the N -electron position vector $\mathbf{R} = (\mathbf{r}_1, \mathbf{r}_2, \dots, \mathbf{r}_N)$. Physically this can be interpreted as the probability that one electron is in region R_1 and that a second electron is in region R_2 and so on is

$$P(\mathbf{r}_1 \in R_1 | \mathbf{r}_2 \in R_2 | \dots | \mathbf{r}_N \in R_N) = \int_{R_1} d^3\mathbf{r}_1 \int_{R_2} d^3\mathbf{r}_2 \dots \int_{R_N} d^3\mathbf{r}_N |\psi(\mathbf{R})|^2. \quad (3)$$

The many-electron time-independent Schrödinger proves extremely difficult to solve exactly because the solution is a function of $3N$ variables. Approximate solutions can be obtained by making different assumptions of varying complexity and with this complexity usually comes an increase in the computational resources required.

Hartree-Fock method

In quantum mechanics two electrons are fundamentally indistinguishable from one another. This and the fact that they are fermions lead to a requirement that the many-electron wave function is antisymmetric under interchange of particles;

$$\psi(\dots, \mathbf{x}_i, \dots, \mathbf{x}_j, \dots) = -\psi(\dots, \mathbf{x}_j, \dots, \mathbf{x}_i, \dots). \quad (4)$$

Where $\mathbf{x}_i = \{\mathbf{r}_i, s_i\}$ represents the space and spin co-ordinates of the electron. The property in Eq. ?? ensures that any two electrons do not have the same set of quantum numbers and hence obey the Pauli exclusion principle[?].

The Hartree-Fock approximation^{???} is the simplest theory that correctly implements the anti-symmetric properties of the many-electron wave function. The simplest wave function with the required antisymmetry is given by the slater determinant

$$D(\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N) = \frac{1}{\sqrt{N!}} \begin{vmatrix} \psi_1(\mathbf{x}_1) & \psi_1(\mathbf{x}_2) & \dots & \psi_1(\mathbf{x}_N) \\ \psi_2(\mathbf{x}_1) & \psi_2(\mathbf{x}_2) & \dots & \psi_2(\mathbf{x}_N) \\ \vdots & \vdots & \ddots & \vdots \\ \psi_N(\mathbf{x}_1) & \psi_N(\mathbf{x}_2) & \dots & \psi_N(\mathbf{x}_N) \end{vmatrix}. \quad (5)$$

The simple antisymmetric slater determinant, Eq. ??, is a starting point for finding the ground state. It can be used as a variational trial function and optimised by minimising the expectation value of the hamiltonian, H with respect to each of the orbitals $\psi_i(\mathbf{r})$. This leads to a set of N self-consistent Hartree-Fock equations describing the single electron orbitals. The solutions to these behave as though each electron is subjected to the mean-field created by all of the other electrons.

This single determinant method includes the effects of the antisymmetry property of the many-electron wave function, Eq. ??, but neglects the electronic correlation arising from the electron-electron Coulomb repulsion. Correlation energies are only a small fraction of the total energy but are still important when considering binding energies[?]. This correlation energy can be accounted for by using a linear combination of slater determinants. However, this leads to another problem as a very large number of determinants are required to accurately describe a many-electron wave function. For a full configuration-interaction calculation the number of determinants required is seen to scale exponentially with system size[?]. This is a big problem for computational resources and other, more efficient methods must be considered to calculate ground state wave functions and energies.ⁱ

ⁱOne such popular method is Density Functional Theory (DFT). This has favourable $N^2 - N^3$ scaling, but often fails for strongly correlated systems and non-local correlation phenomena

Diffusion Monte Carlo and the Fermion Sign Problem

A starting point for all “projector” quantum Monte Carlo methods[?] is to transform the time-dependent Schrödinger equation into imaginary time, $\tau = it$. The so-called imaginary-time Schrödinger equation is then,

$$\frac{\partial \psi}{\partial \tau} = -H\psi, \quad (6)$$

where ψ is the time-dependent many-electron wave function. A formal solution to the imaginary-time Schrödinger equation can be written,

$$\psi(\tau_1 + \delta\tau) = e^{-H\delta\tau} \psi(\tau_1). \quad (7)$$

The state ψ evolves in imaginary time over a duration of $\delta\tau$. If the initial state, $\psi(\tau_1)$, is expanded as a linear combination of ordered energy eigenstates, ϕ_i , with $\epsilon_0 < \epsilon_1 < \dots < \epsilon_\infty$, then

$$\psi(\delta\tau) = \sum_{i=0}^{\infty} c_i e^{-\epsilon_i \delta\tau} \phi_i. \quad (8)$$

From this it is then obvious that any initial state, ψ , that is not orthogonal to the ground-state, ϕ_0 , will evolve to the ground state in the long-time limit,

$$\lim_{\tau \rightarrow \infty} \psi = c_0 e^{-\epsilon_0 \delta\tau} \phi_0. \quad (9)$$

The ground-state is therefore projected out. The aim of all Monte Carlo projector methods, such as Diffusion Monte Carlo (DMC)[?], is to perform a stochastic long-time integration of Eq. ?? in order to retrieve the ground-state.

If the hamiltonian in Eq. ?? is expanded in terms of the kinetic energy and potential termsⁱ, the imaginary-time Schrödinger equation takes on a form that such as van der Waals interactions[?].

ⁱAn adjustable constant energy offset, E_T , is introduced in order to keep Eq. ?? finite.

is analogous to the diffusion equation:

$$-\frac{\partial}{\partial \tau} \psi(\mathbf{R}, \tau) = \left[-\frac{1}{2} \sum_{i=1}^N \nabla_{\mathbf{r}_i}^2 + V(\mathbf{R}) - E_T \right] \psi(\mathbf{R}, \tau). \quad (10)$$

Here, $\psi(\mathbf{R}, \tau)$ may be interpreted as the density of diffusing particles and $(V(\mathbf{R}) - E_T)$ as a rate term describing a potential-dependent change in the particle density. In the Monte Carlo literature these particles are known as “walkers”.




Figure 1: The diffusion of walkers through \mathbf{R} -space. An initially even distribution of walkers is propagated through imaginary time, τ . Under the influence of the potential $V(\mathbf{R})$, the distribution eventually evolves into what is representative of the ground state many-electron wave function $\psi(\mathbf{R})$.

It is well known that the only major obstacle preventing the exact solution of the many-electron Schrödinger equation via stochastic methods, such as DMC, is the fermion sign problem. This problem arises from the antisymmetry of many-electron wave functions, Eq. ?? . As the Schrödinger equation becomes a diffusion equation in imaginary time, Eq. ?? , its lowest energy solution is generally symmetric and nodeless, which directly disagrees with fermion antisymmetry. Stochastic propagation leads directly to this undesired solution in a process known as the “boson catastrophe” ? .

One way of trying to prevent the boson catastrophe is the fixed-node approxi-

mation^{??}. In this, a trial function is used to predict positive and negative regions of the wave function, thus defining a nodal surface. If a walker attempts to cross a node then it is rejected, hence constraining the propagation to disjoint areas of propagation. This procedure would be exact if the applied nodal boundaries coincided with the exact nodal hypersurface of the ground-state electronic wave function[?]. It has in fact proven very difficult to determine accurate fixed-node surfaces for large systems.

Full Configuration-Interaction Quantum Monte Carlo

In 2009 Booth, Alavi and Thom[?] introduced a new quantum Monte Carlo method for the simulation of correlated many-electron systems in full configuration-interaction spaces. The new method is designed to simulate the many-electron imaginary-time Schrödinger using a set of walkers that inhabit Slater determinant space, $\{D_i\}$, and evolve under the guidance of a simple set of rules; spawning, cloning, death and annihilation. The inventors of FCIQMC note three key elements of the algorithm;

1. As with DMC a long-time integration of the imaginary-time Schrödinger equation is performed, however it is achieved in a space of Slater determinants
2. In a similar way to DMC the instantaneous wave function is represented by walkers instead of amplitude coefficients, allowing the description of the FCI wave function stochastically, without storing all amplitudes simultaneously
3. Walkers carry a positive or negative sign which allows for annihilation, where two walkers of opposite sign are destroyed if they coincide on the same determinant

In order to satisfy these three key elements a long-time integration must be performed on an analogue of Eq. ?? that is expressed in a Slater determinant

basis. A set of coupled first-order differential equations for the FCI coefficients are derived[?] :

$$-\frac{dC_i}{d\tau} = (K_{ii} - S)C_i + \sum_{j \neq i} K_{ij}C_j, \quad (11)$$

where

$$K_{ij} \equiv \langle D_i | K | D_j \rangle = \langle D_i | H | D_j \rangle - E_{HF}\delta_{ij}, \quad (12)$$

where S is an adjustable energy offsetⁱ, C_j is the coefficient of the j th determinant in the Slater determinant expansion of the wave function (CI coefficients) and E_{HF} is the Hartree-Fock energy.

The population dynamics algorithm^{???} simulates the set of differential equations, Eq. ???. The algorithm consists of three steps, which are performed at each time step of length $\delta\tau$:

1. Spawning: Each walker located on D_i attempts to spawn another walker on a connected determinant D_j with probability

$$P_s(D_j|D_i) = \frac{\delta\tau|K_{ij}|}{P_{gen}(D_j|D_i)}, \quad (13)$$

where $P_{gen}(D_j|D_i)$ is the calculated probability of generating D_j ⁱⁱ. The sign of the spawned walker is the same as the parent if $K_{ij} < 0$ and is opposite otherwise.

2. Diagonal death/cloning: For each parent (pre-existing walker) compute

$$p_d(D_i) = \delta\tau(K_{ii} - S). \quad (14)$$

If $p_d > 0$ the walker dies (immediately removed from simulation) with probability p_d otherwise it is cloned with probability $|p_d|$.

3. Annihilation: run over all remaining walkers and annihilate all opposite-sign pairs that occupy the same point in antisymmetric Slater determinant

ⁱSimilar to in DMC this energy offset is introduced to keep the population under control

ⁱⁱBooth *et al.*[?] describe how to calculate P_{gen}

space.

Such an algorithm converges on the exact fermionic ground-state of the Hamiltonian in the Slater determinant basis[?]. So in this way, it prevents convergence to a bosonic solution, which is a major problem in other QMC methods. Moreover, FCIQMC simulations do not require any *a priori* information and hence removes the hassle of trying to calculate a nodal structure as is inherent in DMC.

Appendix B: C_h spline derivative estimator

This appendix includes more detail on the calculation of the heat capacity in a uniform magnetic field using the derivative of the DMQMC finite-temperature energy estimator.

Smoothing spline fit

In Sec. ?? it was found that the heat capacity in a uniform magnetic field could be expressed as

$$C_h = -\beta^2 \frac{d\langle H \rangle}{d\beta}. \quad (15)$$

The problem here is that the $\langle H \rangle$ data contains statistical noise and it thus proves difficult to obtain an accurate approximation of the gradient function. However, it is possible to fit a smooth curve to the data using a smoothing spline[?]. A spline is essentially a piece-wise polynomial. The data is split into regions and each region is approximated by a polynomial, usually a cubic.

For a sequence of measurements M_i that can be modelled as $M_i = F(x_i)$, where $\{x_i | x_1 < x_2 < \dots < x_N, x_i \in \mathcal{Z}\}$ and F is some function, the smoothing

spline estimate \bar{F} of F is defined as the function that minimises

$$\sum_{i=1}^n (M_i - \bar{F}(x_i))^2 + \eta \int_{x_1}^{x_N} \bar{F}''(x)^2 dx. \quad (16)$$

Here $\eta \geq 0$ is the smoothing parameter which controls the smoothing characteristics \bar{F} . For $\eta \rightarrow 0$ \bar{F} becomes an interpolating spline and for $\eta \rightarrow \infty$ \bar{F} converges to a linear least squares estimate.

Python has a smoothing spline implementation, `splev`, in the `scipy.interpolate` sub-package. In this the smoothing parameter is chosen so that it is between $N - \sqrt{2N} < \eta < N + \sqrt{2N}$ and the best fit of the data can be found by trial and error. The `splev` function can also return derivatives of \bar{F} , which proved useful for calculating the C_h .

So during this investigation the DMQMC spline derivative estimate of C_h was determined in a post-processing step. This step was inelegant as it relied on some trial and error and it was difficult to obtain errors.

Appendix C: Code

This appendix presents some of the main components of the DMQMC implementation in FORTRAN 90 that were developed during this project. The `dmqmc` module in Code 1 contains the main outline of the algorithm as presented in Sec. ?? . `dmqmc_procedures.f90` and `dmqmc_estimators.f90` in Code 2 and Code 3 contain some important general procedures and the implementation of the estimators respectively.

Code 1: `dmqmc.f90`

```
1  module dmqmc
2
3  ! Main loop for performing DMQMC calculations, similar to the
4  ! file fciqmc.f90 which carries out the main loop in FCIQMC
5  ! calculations.
6
7  use fciqmc_data
8  use proc_pointers
9  implicit none
10
11 contains
12
13  subroutine do_dmqmc()
14
15      ! Run DMQMC calculation. We run from a beta=0 to a value of beta
16      ! specified by the user and then repeat this main loop beta_loops
17      ! times, to accumulate statistics for each value for beta.
18
19      use parallel
20      use annihilation, only: direct_annihilation
21      use basis, only: basis_length, bit_lookup, nbasis
```

```

22     use death, only: stochastic_death
23     use determinants, only: det_info, alloc_det_info, dealloc_det_info
24     use dmqmc_procedures, only: random_distribution_heisenberg
25     use dmqmc_procedures, only: update_sampling_weights, output_and_alter_weights
26     use dmqmc_estimators, only: update_dmqmc_estimators, call_dmqmc_estimators
27     use dmqmc_estimators, only: call_rdm_procedures
28     use excitations, only: excit, get_excitation_level
29     use fciqmc_common
30     use fciqmc_restart, only: dump_restart
31     use interact, only: fciqmc_interact
32     use system, only: nel
33     use calc, only: seed, doing_dmqmc_calc, dmqmc_energy
34     use calc, only: dmqmc_staggered_magnetisation, dmqmc_energy_squared
35     use dSFMT_interface, only: dSFMT_init
36     use utils, only: int_fmt
37
38     integer :: idet, ireport, icycle, iparticle, iteration
39     integer :: beta_cycle
40     integer(lint) :: nparticles_old(sampling_size)
41     integer(lint) :: n attempts, nparticles_start_report
42     type(det_info) :: cdet1, cdet2
43     integer :: nspawned, ndeath
44     type(excit) :: connection
45     integer :: spawning_end
46     logical :: soft_exit
47     real :: t1, t2
48
49     ! Allocate det_info components. We need two cdet objects
50     ! for each 'end' which may be spawned from in the DMQMC algorithm.
51     call alloc_det_info(cdet1, .false.)
52     call alloc_det_info(cdet2, .false.)
53
54     ! Main DMQMC loop.
55     if (parent) call write_fciqmc_report_header()
56     ! Initialise timer.
57     call cpu_time(t1)
58
59     initial_shift = shift
60     ! When we accumulate data throughout a run, we are actually accumulating
61     ! results from the psips distribution from the previous iteration.
62     ! For example, in the first iteration, the trace calculated will be that
63     ! of the initial distribution, which corresponds to beta=0. Hence, in the
64     ! output we subtract one from the iteration number, and run for one more
65     ! report loop, as implemented in the line of code below.

```

```

66     nreport = nreport+1
67
68     do beta_cycle = 1, beta_loops
69         ! Reset the current position in the spawning array to be the
70         ! slot preceding the first slot.
71         spawning_head = spawning_block_start
72         tot_walkers = 0
73         shift = initial_shift
74         nparticles = 0
75         if (allocated(reduced_density_matrix)) reduced_density_matrix = 0
76         if (dmqmc_vary_weights) dmqmc_accumulated_probs = 1.0_p
77         if (dmqmc_find_weights) excit_distribution = 0
78         vary_shift = .false.
79
80         ! Need to place psips randomly along the diagonal at the
81         ! start of every iteration. Pick orbitals randomly, each
82         ! with equal probability, so that when electrons are placed
83         ! on these orbitals they will have the correct spin and symmetry.
84         call dmqmc_initial_distribution_ptr()
85
86         call direct_annihilation()
87
88         if (beta_cycle .ne. 1) then
89             ! Reset the random number generator with seed = seed + nprocs
90             seed = seed + nprocs
91             call dSFMT_init(seed + iproc)
92             if (parent) then
93                 write (6,'(a32,i7)') &
94                     " # Resetting beta... Beta loop =", beta_cycle
95                 write (6,'(a52,'//int_fmt(seed,1)//',a1)') &
96                     " # Resetting random number generator with a seed of:", seed, "."
97             end if
98         end if
99
100        nparticles_old = nint(D0_population)
101
102        do ireport = 1, nreport
103            ! Zero report cycle quantities.
104            rspawn = 0.0_p
105            trace = 0
106            estimator_numerators = 0
107            if (calculate_excit_distribution) excit_distribution = 0
108            nparticles_start_report = nparticles_old(1)
109

```

```

110         do icycle = 1, ncycles
111             spawning_head = spawning_block_start
112             iteration = (ireport-1)*ncycles + icycle
113
114             ! Number of spawning attempts that will be made.
115             ! Each particle and each end gets to attempt to
116             ! spawn onto a connected determinant and a chance
117             ! to die/clone.
118             nattempts = 4*nparticles(1)
119
120             ! Reset death counter
121             ndeath = 0
122
123             do idet = 1, tot_walkers ! loop over walkers/dets
124
125                 ! f points to the bitstring that is spawning, f2 to the
126                 ! other bit string.
127                 cdet1%f => walker_dets(:basis_length,idet)
128                 cdet1%f2 => walker_dets((basis_length+1):(2*basis_length),idet)
129                 cdet2%f => walker_dets((basis_length+1):(2*basis_length),idet)
130                 cdet2%f2 => walker_dets(:basis_length,idet)
131
132                 ! Decode and store the the relevant information for
133                 ! both bitstrings. Both of these bitstrings are required
134                 ! to refer to the correct element in the density matrix.
135                 call decoder_ptr(cdet1%f, cdet1)
136                 call decoder_ptr(cdet2%f, cdet2)
137
138                 ! Call wrapper function which calls all requested estimators
139                 ! to be updated, and also always updates the trace separately.
140
141                 if (icycle == 1) call call_dmcmc_estimators(idet, iteration)
142
143                 do iparticle = 1, abs(walker_population(1,idet))
144                     ! Spawn from the first end.
145                     spawning_end = 1
146                     ! Attempt to spawn.
147                     call spawner_ptr(cdet1, walker_population(1,idet), nspawned, connection)
148                     ! Spawn if attempt was successful.
149                     if (nspawned /= 0) then
150                         call create_spawned_particle_dm_ptr(cdet1%f, cdet2%f, connection, nspawned,
151                             spawning_end)
152                     end if

```

```

153         ! Now attempt to spawn from the second end.
154         spawning_end = 2
155         call spawner_ptr(cdet2, walker_population(1,idet), nspawned, connection)
156         if (nspawned /= 0) then
157             call create_spawned_particle_dm_ptr(cdet2%f, cdet1%f, connection, nspawned,
158                 spawning_end)
159         end if
160     end do
161
162     ! Clone or die.
163     ! We have contirbutions to the clone/death step from both ends of the
164     ! current walker. We do both of these at once by using walker_data(1,idet)
165     ! which, when running a DMQMC algorithm, stores the average of the two diagonal
166     ! elements corresponding to the two indicies of the density matrix (the two ends).
167     call stochastic_death(walker_data(1,idet), walker_population(1,idet), nparticles(1),
168         ndeath)
169 end do
170
171 ! Add the spawning rate (for the processor) to the running
172 ! total.
173 rspawn = rspawn + spawning_rate(ndeath, nattempts)
174
175 ! Perform the annihilation step where the spawned walker list is merged with
176 ! the main walker list, and walkers of opposite sign on the same sites are
177 ! annihilated.
178 call direct_annihilation()
179
180 ! If doing importance sampling *and* varying the weights of the trial function, call a
181 ! routine
182 ! to update these weights and alter the number of psips on each excitation level accordingly
183 .
184 if (dmqmc_vary_weights .and. iteration <= finish_varying_weights) call
185     update_sampling_weights()
186
187 end do
188
189 ! If averaging the shift to use in future beta loops, add contirubtion from this report.
190 if (average_shift_until > 0) shift_profile(ireport) = shift_profile(ireport) + shift
191
192 ! Update the shift and desired thermal quantites.
193 call update_dmqmc_estimators(nparticles_old, ireport)
194
195 call cpu_time(t2)
196

```

```

192         ! t1 was the time at the previous iteration, t2 the current time.
193         ! t2-t1 is thus the time taken by this report loop.
194         if (parent) call write_fciqmc_report(ireport, nparticles_start_report, t2-t1)
195
196         ! cpu_time outputs an elapsed time, so update the reference timer.
197         t1 = t2
198
199         call fciqmc_interact(soft_exit)
200         if (soft_exit) exit
201
202     end do
203
204     if (soft_exit) exit
205
206     ! If have just finished last beta loop of accumulating the shift, then perform
207     ! the averaging and set average_shift_until to -1. This tells the shift update
208     ! algorithm to use the values for shift stored in shift_profile.
209     if (beta_cycle == average_shift_until) then
210         shift_profile = shift_profile/average_shift_until
211         average_shift_until = -1
212     end if
213
214     ! Calculate and output all requested estimators based on the reduced density matrix.
215     if (doing_reduced_dm) call call_rdm_procedures()
216     ! Calculate and output new weights based on the psip distribution in the previous loop.
217     if (dmqmc_find_weights) call output_and_alter_weights()
218 end do
219
220 if (parent) then
221     call write_fciqmc_final(ireport)
222     write (6, '( )')
223 end if
224
225 call load_balancing_report()
226
227 if (dump_restart_file) call dump_restart(mc_cycles_done+ncycles*nreport, nparticles_old(1))
228
229 call dealloc_det_info(cdet1, .false.)
230 call dealloc_det_info(cdet2, .false.)
231
232 end subroutine do_dmqmc
233
234 end module dmqmc

```

Code 2: dmqmc_procedures.f90

```
1 module dmqmc_procedures
2
3 use const
4 implicit none
5
6 contains
7
8   subroutine init_dmqmc()
9
10      use basis, only: basis_length, total_basis_length, bit_lookup, basis_lookup
11      use calc, only: doing_dmqmc_calc, dmqmc_calc_type, dmqmc_energy, dmqmc_energy_squared
12      use calc, only: dmqmc_staggered_magnetisation, dmqmc_correlation
13      use checking, only: check_allocate, check_deallocate
14      use fciqmc_data, only: trace, energy_index, energy_squared_index, correlation_index
15      use fciqmc_data, only: staggered_mag_index, estimator_numerators, subsystem_A_size
16      use fciqmc_data, only: subsystem_A_mask, subsystem_B_mask, subsystem_A_bit_positions
17      use fciqmc_data, only: subsystem_A_list, dmqmc_factor, number_dmqmc_estimators, ncycles
18      use fciqmc_data, only: reduced_density_matrix, doing_reduced_dm, tau, dmqmc_weighted_sampling
19      use fciqmc_data, only: correlation_mask, correlation_sites, half_density_matrix
20      use fciqmc_data, only: dmqmc_sampling_probs, dmqmc_accumulated_probs, flip_spin_matrix
21      use fciqmc_data, only: doing_concurrence, calculate_excit_distribution, excit_distribution
22      use fciqmc_data, only: nreport, average_shift_until, shift_profile, dmqmc_vary_weights
23      use fciqmc_data, only: finish_varying_weights, weight_altering_factors, dmqmc_find_weights
24      use parallel, only: parent
25      use system, only: system_type, heisenberg, nsites, max_number_excitations
26
27      integer :: ierr
28      integer :: i, ipos, basis_find, bit_position, bit_element
29
30      number_dmqmc_estimators = 0
31      trace = 0
32
33      if (doing_dmqmc_calc(dmqmc_energy)) then
34         number_dmqmc_estimators = number_dmqmc_estimators + 1
35         energy_index = number_dmqmc_estimators
36      end if
37      if (doing_dmqmc_calc(dmqmc_energy_squared)) then
38         number_dmqmc_estimators = number_dmqmc_estimators + 1
39         energy_squared_index = number_dmqmc_estimators
40      end if
41      if (doing_dmqmc_calc(dmqmc_correlation)) then
42         number_dmqmc_estimators = number_dmqmc_estimators + 1
43         correlation_index = number_dmqmc_estimators
```

```

44     allocate(correlation_mask(1:basis_length), stat=ierr)
45     call check_allocate('correlation_mask',basis_length,ierr)
46     correlation_mask = 0
47     do i = 1, 2
48         bit_position = bit_lookup(1,correlation_sites(i))
49         bit_element = bit_lookup(2,correlation_sites(i))
50         correlation_mask(bit_element) = ibset(correlation_mask(bit_element), bit_position)
51     end do
52 end if
53 if (doing_dmcmc_calc(dmcmc_staggered_magnetisation)) then
54     number_dmcmc_estimators = number_dmcmc_estimators + 1
55     staggered_mag_index = number_dmcmc_estimators
56 end if
57
58 allocate(estimator_numerators(1:number_dmcmc_estimators), stat=ierr)
59 call check_allocate('estimator_numerators',number_dmcmc_estimators,ierr)
60 estimator_numerators = 0
61
62 if (calculate_excit_distribution .or. dmcmc_find_weights) then
63     allocate(excit_distribution(0:max_number_excitations), stat=ierr)
64     call check_allocate('excit_distribution',max_number_excitations+1,ierr)
65     excit_distribution = 0.0_p
66 end if
67
68 if (average_shift_until > 0) then
69     allocate(shift_profile(1:nreport+1), stat=ierr)
70     call check_allocate('shift_profile',nreport+1,ierr)
71     shift_profile = 0.0_p
72 end if
73
74 ! In DMQMC we want the spawning probabilities to have an extra factor of a half,
75 ! because we spawn from two different ends with half probability. To avoid having
76 ! to multiply by an extra variable in every spawning routine to account for this, we
77 ! multiply the time step by 0.5 instead, then correct this in the death step (see below).
78 tau = tau*0.5_p
79
80 ! Set dmcmc_factor to 2 so that when probabilities in death.f90 are multiplied
81 ! by this factor it cancels the factor of 0.5 introduced into the timestep in DMQMC.
82 ! Every system uses the same death routine, so this factor only needs to be added once.
83 ! This factor is also used in updated the shift, where the true tau is needed.
84 dmcmc_factor = 2.0_p
85
86 if (dmcmc_weighted_sampling) then
87     ! dmcmc_sampling_probs stores the factors by which probabilities are to
88     ! be reduced when spawning away from the diagonal. The trial function required

```

```

88      ! from these probabilities, for use in importance sampling, is actually that of
89      ! the accumulated factors, ie, if dmqmc_sampling_probs = (a, b, c, ...) then
90      ! dmqmc_accumulated_factors = (1, a, ab, abc, ...). This is the array which we
91      ! need to create and store. dmqmc_sampling_probs is no longer needed and so can
92      ! be deallocated. Also, the user may have only input factors for the first few
93      ! excitation levels, but we need to store factors for all levels, as done below.
94      if (.not.allocated(dmqmc_sampling_probs)) then
95          allocate(dmqmc_sampling_probs(1:max_number_excitations), stat=ierr)
96          call check_allocate('dmqmc_sampling_probs',max_number_excitations,ierr)
97          dmqmc_sampling_probs = 1.0_p
98      end if
99      if (half_density_matrix) dmqmc_sampling_probs(1) = dmqmc_sampling_probs(1)*2.0_p
100     allocate(dmqmc_accumulated_probs(0:max_number_excitations), stat=ierr)
101     call check_allocate('dmqmc_accumulated_probs',max_number_excitations+1,ierr)
102     dmqmc_accumulated_probs(0) = 1.0_p
103     do i = 1, size(dmqmc_sampling_probs)
104         dmqmc_accumulated_probs(i) = dmqmc_accumulated_probs(i-1)*dmqmc_sampling_probs(i)
105     end do
106     dmqmc_accumulated_probs(size(dmqmc_sampling_probs)+1:max_number_excitations) = &
107         dmqmc_accumulated_probs(size(dmqmc_sampling_probs))
108     if (dmqmc_vary_weights) then
109         ! Allocate an array to store the factors by which the weights will change each
110         ! iteration.
111         allocate(weight_altering_factors(0:max_number_excitations), stat=ierr)
112         call check_allocate('weight_altering_factors',max_number_excitations+1,ierr)
113         weight_altering_factors = dble(dmqmc_accumulated_probs)*(1/dble(finish_varying_weights))
114         ! If varying the weights, start the accumulated probabilities as all 1.0
115         ! initially, and then alter them gradually later.
116         dmqmc_accumulated_probs = 1.0_p
117     end if
118 else
119     allocate(dmqmc_accumulated_probs(0:max_number_excitations), stat=ierr)
120     call check_allocate('dmqmc_accumulated_probs',max_number_excitations+1,ierr)
121     dmqmc_accumulated_probs = 1.0_p
122     if (half_density_matrix) dmqmc_accumulated_probs(1:max_number_excitations) &
123         = 2.0_p*dmqmc_accumulated_probs(1:max_number_excitations)
124 end if
125
126 ! If doing a reduced density matrix calculation, then allocate and define the
127 ! bit masks that have 1's at the positions referring to either subsystems A or B.
128 if (doing_reduced_dm) then
129     subsystem_A_size = ubound(subsystem_A_list,1)
130     allocate(subsystem_A_mask(1:basis_length), stat=ierr)
131     call check_allocate('subsystem_A_mask',basis_length,ierr)

```

```

132     allocate(subsystem_B_mask(1:basis_length), stat=ierr)
133     call check_allocate('subsystem_B_mask',basis_length,ierr)
134     allocate(subsystem_A_bit_positions(subsystem_A_size,2), stat=ierr)
135     call check_allocate('subsystem_A_bit_positions',2*subsystem_A_size,ierr)
136     subsystem_A_mask = 0
137     subsystem_B_mask = 0
138     subsystem_A_bit_positions = 0
139     ! For the Heisenberg model only currently.
140     if (system_type==heisenberg) then
141         do i = 1, subsystem_A_size
142             bit_position = bit_lookup(1,subsystem_A_list(i))
143             bit_element = bit_lookup(2,subsystem_A_list(i))
144             subsystem_A_mask(bit_element) = ibset(subsystem_A_mask(bit_element), bit_position)
145             subsystem_A_bit_positions(i,1) = bit_position
146             subsystem_A_bit_positions(i,2) = bit_element
147         end do
148         subsystem_B_mask = subsystem_A_mask
149         ! We cannot just flip the mask for system A to get that for system B, because
150         ! there the trailing bits on the end don't refer to anything and should be
151         ! set to 0. So, first set these to 1 and then flip all the bits.
152         do ipos = 0, i0_end
153             basis_find = basis_lookup(ipos, basis_length)
154             if (basis_find == 0) then
155                 subsystem_B_mask(basis_length) = ibset(subsystem_B_mask(basis_length),ipos)
156             end if
157         end do
158         subsystem_B_mask = not(subsystem_B_mask)
159     end if
160     if (subsystem_A_size <= int(nsites/2)) then
161         ! In this case, for an ms = 0 subspace (as the ground state of the Heisenberg model
162         ! will be) then any combination of spins can occur in the subsystem, from all spins
163         ! down to all spins up. Hence the total size of the reduced density matrix will be
164         ! 2**(number of spins in subsystem A).
165         allocate(reduced_density_matrix(2**subsystem_A_size,2**subsystem_A_size), stat=ierr)
166         call check_allocate('reduced_density_matrix', 2**(2*subsystem_A_size),ierr)
167         reduced_density_matrix = 0
168     else if (subsystem_A_size == nsites) then
169         allocate(reduced_density_matrix(2**subsystem_A_size,2**subsystem_A_size), stat=ierr)
170         call check_allocate('reduced_density_matrix', 2**(2*subsystem_A_size),ierr)
171         reduced_density_matrix = 0
172     end if
173 end if
174
175 ! If doing concurrence calculation then construct and store the 4x4 flip spin matrix i.e.

```

```

176         ! \sigma_y \otimes \sigma_y
177
178         if (doing_concurrence) then
179             allocate(flip_spin_matrix(4,4), stat=ierr)
180             call check_allocate('flip_spin_matrix', 16,ierr)
181             flip_spin_matrix = 0._p
182             flip_spin_matrix(1,4) = -1._p
183             flip_spin_matrix(4,1) = -1._p
184             flip_spin_matrix(3,2) = 1._p
185             flip_spin_matrix(2,3) = 1._p
186         end if
187
188     end subroutine init_dmcmc
189
190     subroutine random_distribution_heisenberg()
191
192         ! For the Heisenberg model only. Distribute the initial number of psips
193         ! along the main diagonal. Each diagonal element should be chosen
194         ! with the same probability.
195
196         ! Currently this creates psips with Ms = ms_in only.
197
198         ! If we have number of sites = nsites,
199         ! and total spin value = ms_in,
200         ! then number of up spins is equal to up_spins = (ms_in + nsites)/2.
201
202         ! Start from state with all spins down, then choose the above number of
203         ! spins to flip up with equal probability.
204
205         use basis, only: nbasis, basis_length, bit_lookup
206         use calc, only: ms_in
207         use dSFMT_interface, only: genrand_real2
208         use fciqmc_data, only: D0_population
209         use parallel
210         use system, only: nsites
211
212         integer :: i, up_spins, rand_basis, bits_set
213         integer :: bit_element, bit_position, npsips
214         integer(i0) :: f(basis_length)
215         real(dp) :: rand_num
216
217         up_spins = (ms_in+nsites)/2
218         npsips = int(D0_population/nprocs)
219         ! If initial number of psips does not split evenly between all processors,

```

```

220      ! add the leftover psips to the first processors in order.
221      if (D0_population-(nprocs*int(D0_population/nprocs)) > iproc) npsips = npsips+1
222
223      do i = 1, npsips
224
225          ! Start with all spins down.
226          f = 0
227          bits_set = 0
228
229          do
230              ! If half the spins are now flipped up, we have our basis
231              ! function fully created, so exit the loop.
232              if (bits_set==up_spins) exit
233              ! Choose a random spin to flip.
234              rand_num = genrand_real2()
235              rand_basis = ceiling(rand_num*nbasis)
236              ! Find the corresponding positions for this spin.
237              bit_position = bit_lookup(1,rand_basis)
238              bit_element = bit_lookup(2,rand_basis)
239              if (.not. btest(f(bit_element),bit_position)) then
240                  ! If not flipped up, flip the spin up.
241                  f(bit_element) = ibset(f(bit_element),bit_position)
242                  bits_set = bits_set + 1
243              end if
244          end do
245
246          ! Now call a routine to add the corresponding diagonal element to
247          ! the spawned walkers list.
248          call create_particle(f,f,1)
249
250      end do
251
252  end subroutine random_distribution_heisenberg
253
254  subroutine create_particle(f1,f2,nspawn)
255
256      ! Create a psip on a diagonal element of the density
257      ! matrix by adding it to the spawned walkers list. This
258      ! list can then be sorted correctly by the direct_annihilation
259      ! routine
260
261      ! In:
262      ! f_new: Bit string representation of index of the diagonal
263      ! element upon which a new psip shall be placed.

```

```

264
265     use hashing
266     use basis, only: basis_length, total_basis_length
267     use fciqmc_data, only: spawned_walkers, spawning_head, spawned_pop
268     use parallel
269
270     integer(i0), intent(in) :: f1(basis_length), f2(basis_length)
271     integer, intent(in) :: nspawn
272     integer(i0) :: f_new(total_basis_length)
273 #ifndef PARALLEL
274     integer, parameter :: iproc_spawn = 0
275 #else
276     integer :: iproc_spawn
277 #endif
278
279     ! Create the bitstring of the psip.
280     f_new = 0
281     f_new(:basis_length) = f1
282     f_new((basis_length+1):(total_basis_length)) = f2
283
284 #ifdef PARALLEL
285     ! Need to determine which processor the spawned walker should be sent to.
286     iproc_spawn = modulo(murmurhash_bit_string(f_new, &
287                                     (total_basis_length)), nprocs)
288 #endif
289
290     ! Move to the next position in the spawning array.
291     spawning_head(iproc_spawn) = spawning_head(iproc_spawn) + 1
292
293     ! Set info in spawning array.
294     ! Zero it as not all fields are set.
295     spawned_walkers(:,spawning_head(iproc_spawn)) = 0
296     ! indices 1 to total_basis_length store the bitstring.
297     spawned_walkers((2*basis_length),spawning_head(iproc_spawn)) = f_new
298     ! The final index stores the number of psips created.
299     spawned_walkers((2*basis_length)+1,spawning_head(iproc_spawn)) = nspawn
300
301 end subroutine create_particle
302
303 subroutine decode_dm_bitstring(f, index1, index2)
304
305     ! This function maps an input DMQMC bitstring to two indices
306     ! giving the corresponding position of the bitstring in the reduced
307     ! density matrix.

```

```

308
309     use basis, only: total_basis_length, basis_length
310     use fciqmc_data, only: subsystem_A_bit_positions, subsystem_A_bit_positions
311     use fciqmc_data, only: subsystem_A_size
312
313     integer(i0), intent(in) :: f(total_basis_length)
314     integer(i0), intent(out) :: index1, index2
315     integer :: i
316
317     ! Start from all bits down, so that we can flip bits up one by one.
318     index1 = 0
319     index2 = 0
320
321     ! Loop over all the sites in the sublattice considered for the reduced density matrix.
322     do i = 1, subsystem_A_size
323         ! If the spin is up, flip the corresponding bit in the first index up.
324         if (btest(f(subsystem_A_bit_positions(i,2)),subsystem_A_bit_positions(i,1))) &
325             index1 = ibset(index1,i-1)
326         ! Similarly for the second index, by looking at the second end of the bitstring.
327         if (btest(f(subsystem_A_bit_positions(i,2)+basis_length),subsystem_A_bit_positions(i,1))) &
328             index2 = ibset(index2,i-1)
329     end do
330
331     ! The process above maps to numbers between 0 and 2^subsystem_A_size-1, but the smallest
332     ! and largest of the reduced density matrix are one more than these, so add one...
333     index1 = index1+1
334     index2 = index2+1
335
336 end subroutine decode_dm_bitstring
337
338 subroutine update_sampling_weights()
339
340     ! This routine updates the values of the weights used in importance sampling. It also
341     ! removes or adds psips from the various levels accordingly.
342
343     ! In:
344     ! iteration: The current iteration in the beta loop.
345
346     use annihilation, only: remove_unoccupied_dets
347     use basis, only: basis_length, total_basis_length
348     use excitations, only: get_excitation_level
349     use fciqmc_data, only: dmqmc_accumulated_probs, finish_varying_weights
350     use fciqmc_data, only: weight_altering_factors, tot_walkers, walker_dets, walker_population
351     use fciqmc_data, only: nparticles

```

```

352     use dSFMT_interface, only: genrand_real2
353
354     integer :: idet, excit_level, nspawn, sign_factor, old_population
355     real(p) :: new_factor
356     real(dp) :: rand_num, prob
357
358     ! Alter weights for the next iteration.
359     dmcmc_accumulated_probs = dble(dmcmc_accumulated_probs)*weight_altering_factors
360
361     ! When the weights for an excitation level are increased by a factor, the number
362     ! of psips on that level has to decrease by the same factor, else the wavefunction
363     ! which the psips represent will not be the correct importance sampled wavefunction
364     ! for the new weights. The code below loops over every psips and destorys (or creates)
365     ! it with the appropriate probability.
366     do idet = 1, tot_walkers
367         excit_level = get_excitation_level(walker_dets(1:basis_length,idet), &
368             walker_dets(basis_length+1:total_basis_length,idet))
369         old_population = abs(walker_population(1,idet))
370         rand_num = genrand_real2()
371         ! If weight_altering_factors(excit_level) > 1, need to kill psips.
372         ! If weight_altering_factors(excit_level) < 1, need to create psips.
373         prob = abs(1.0_dp - weight_altering_factors(excit_level)**(-1))*old_population
374         nspawn = int(prob)
375         prob = prob - nspawn
376         if (rand_num < prob) nspawn = nspawn + 1
377         if (weight_altering_factors(excit_level) > 1.0_dp) then
378             sign_factor = -1
379         else
380             sign_factor = +1
381         end if
382         nspawn = sign(nspawn,walker_population(1,idet)*sign_factor)
383         ! Update the population on this determinant.
384         walker_population(1,idet) = walker_population(1,idet) + nspawn
385         ! Update the total number of walkers
386         nparticles(1) = nparticles(1) - old_population + abs(walker_population(1,idet))
387     end do
388
389     ! Call the annihilation routine to update the main walker list, as some
390     ! sites will now have no psips on and so need removing from the simulation.
391     call remove_unoccupied_dets()
392
393 end subroutine update_sampling_weights
394
395 subroutine output_and_alter_weights()

```

```

396
397     ! This routine will alter and output the sampling weights used in importance
398     ! sampling. It uses the excitation distribution, calculated on the beta loop
399     ! which has just finished, and finds the weights needed so that each excitation
400     ! level will have roughly equal numbers of psips in the next loop. For example,
401     ! to find the weights of psips on the 1st excitation level, divide the number of
402     ! psips on the 1st excitation level by the number on the 0th level, then multiply
403     ! the old sampling weight by this number to give the new weight. This can be used
404     ! when the weights are being introduced gradually each beta loop, too. The weights
405     ! are output and can then be used in future DMQMC runs.
406
407     use fciqmc_data, only: dmqmc_sampling_probs, dmqmc_accumulated_probs
408     use fciqmc_data, only: excit_distribution, finish_varying_weights
409     use fciqmc_data, only: dmqmc_vary_weights, weight_altering_factors
410     use parallel
411     use system, only: max_number_excitations
412
413     integer :: i, ierr
414     #ifdef PARALLEL
415         real(p) :: merged_excit_dist(max_number_excitations)
416         call mpi_allreduce(excit_distribution, merged_excit_dist, max_number_excitations, &
417             MPI_REAL8, MPI_SUM, MPI_COMM_WORLD, ierr)
418
419         excit_distribution = merged_excit_dist
420     #endif
421
422     ! It is assumed that there is an even maximum number of excitations.
423     do i = 1, (max_number_excitations/2)
424         ! Don't include levels where there are very few psips accumulated.
425         if (excit_distribution(i-1) > 10.0_p .and. excit_distribution(i) > 10.0_p) then
426             ! Alter the sampling weights using the relevant excitation distribution.
427             dmqmc_sampling_probs(i) = dmqmc_sampling_probs(i)*&
428                 (excit_distribution(i)/excit_distribution(i-1))
429             dmqmc_sampling_probs(max_number_excitations+1-i) = dmqmc_sampling_probs(i)**(-1)
430         end if
431     end do
432
433     ! Recalculate dmqmc_accumulated_probs with the new weights.
434     do i = 1, max_number_excitations
435         dmqmc_accumulated_probs(i) = dmqmc_accumulated_probs(i-1)*dmqmc_sampling_probs(i)
436     end do
437
438     ! If dmqmc_vary_weights is true then the weights are to be introduced gradually at the
439     ! start of each beta loop. This required redefining weight_altering_factors to coincide

```

```

440      ! with the new sampling weights.
441      if (dmqmc_vary_weights) then
442          weight_altering_factors = dble(dmqmc_accumulated_probs)**(1/dble(finish_varying_weights))
443          ! Reset the weights for the next loop.
444          dmqmc_accumulated_probs = 1.0_p
445      end if
446
447      if (parent) then
448          ! Print out weights in a form which can be copied into an input file.
449          write(6, '(a31,2X)', advance = 'no') ' # Importance sampling weights:'
450          do i = 1, max_number_excitations
451              write (6, '(es12.4,2X)', advance = 'no') dmqmc_sampling_probs(i)
452          end do
453          write (6, '()', advance = 'yes')
454      end if
455
456      end subroutine output_and_alter_weights
457
458 end module dmqmc_procedures

```

Code 3: dmqmc_estimators.f90

```
1  module dmqmc_estimators
2
3  use const
4
5  implicit none
6
7  contains
8
9      subroutine update_dmqmc_estimators(ntot_particles_old, ireport)
10
11          ! Update the shift and average the shift and estimators
12
13          ! Should be called every report loop in an DMQMC calculation.
14
15          ! Inout:
16          ! ntot_particles_old: total number (across all processors) of
17          ! particles in the simulation at end of the previous report loop.
18          ! Returns the current total number of particles for use in the
19          ! next report loop.
20
21          use calc, only: doing_dmqmc_calc, dmqmc_energy, dmqmc_staggered_magnetisation
22          use calc, only: dmqmc_energy_squared
23          use checking, only: check_allocate
24          use energy_evaluation, only: update_shift
25          use fciqmc_data, only: nparticles, sampling_size, target_particles, rspawn
26          use fciqmc_data, only: shift, vary_shift, nreport
27          use fciqmc_data, only: estimator_numerators, number_dmqmc_estimators
28          use fciqmc_data, only: nreport, ncycles, trace, calculate_excit_distribution
29          use fciqmc_data, only: excit_distribution, average_shift_until, shift_profile
30          use parallel
31
32          integer(lint), intent(inout) :: ntot_particles_old(sampling_size)
33          integer, intent(in) :: ireport
34          integer(lint) :: ntot_particles(sampling_size)
35
36          #ifdef PARALLEL
37              real(dp), allocatable :: ir(:)
38              real(dp), allocatable :: ir_sum(:)
39              integer :: ierr, array_size
40
41              array_size = sampling_size+2+number_dmqmc_estimators
42              if (calculate_excit_distribution) array_size = array_size + size(excit_distribution)
43
```

```

44     allocate(ir(1:array_size), stat=ierr)
45     call check_allocate('ir',array_size,ierr)
46     allocate(ir_sum(1:array_size), stat=ierr)
47     call check_allocate('ir_sum',array_size,ierr)
48
49     ! Need to sum the number of particles and other quantites over all processors.
50     ir(1:sampling_size) = nparticles
51     ir(sampling_size+1) = rspawn
52     ir(sampling_size+2) = trace
53     ir(sampling_size+3:sampling_size+2+number_dmcmc_estimators) = estimator_numerators
54     if (calculate_excit_distribution) ir(sampling_size+3+number_dmcmc_estimators:array_size) =
        excit_distribution
55     ! Merge the lists from each processor together.
56     call mpi_allreduce(ir, ir_sum, size(ir), MPI_REAL8, MPI_SUM, MPI_COMM_WORLD, ierr)
57     ntot_particles = nint(ir_sum(1:sampling_size),lint)
58     rspawn = ir_sum(sampling_size+1)
59     trace = nint(ir_sum(sampling_size+2))
60     estimator_numerators = ir_sum(sampling_size+3:sampling_size+2+number_dmcmc_estimators)
61     if (calculate_excit_distribution) excit_distribution = ir_sum(sampling_size+3+number_dmcmc_estimators:
        array_size)
62 #else
63     ntot_particles = nparticles
64 #endif
65
66     ! If average_shift_until = -1 then it means that the shift should be updated to
67     ! use the values of shift stored in shift_profile. Otherwise, use the standard update
68     ! routine.
69     if (average_shift_until == -1) then
70         if (ireport < nreport) shift = shift_profile(ireport+1)
71     else
72         if (vary_shift) call update_shift(ntot_particles_old(1), ntot_particles(1), ncycles)
73         if (nparticles(1) > target_particles .and. .not.vary_shift) vary_shift = .true.
74     end if
75
76     ntot_particles_old = ntot_particles
77     rspawn = rspawn/(ncycles*nprocs)
78
79 end subroutine update_dmcmc_estimators
80
81 subroutine call_dmcmc_estimators(idet, iteration)
82
83     ! This function calls the processes to update the estimators which
84     ! have been requested by the user to be calculated.
85     ! First, calculate the excitation level between the two bitsrtings

```

```

86      ! corresponding to the the two ends. Then add the contribution from
87      ! the current density matrix element to the trace, which is always
88      ! calculated. Then call other estimators, as required.
89
90      ! In:
91      ! idet: Current position in the main bitstring list.
92      ! iteration: current Monte Carlo cycle.
93
94      use basis, only: basis_length, total_basis_length
95      use calc, only: doing_dmcmc_calc, dmcmc_energy, dmcmc_staggered_magnetisation
96      use calc, only: dmcmc_energy_squared, dmcmc_correlation
97      use excitations, only: get_excitation, excit
98      use fciqmc_data, only: walker_dets, walker_population, trace, doing_reduced_dm
99      use fciqmc_data, only: dmcmc_accumulated_probs, start_averaging, dmcmc_find_weights
100     use fciqmc_data, only: calculate_excit_distribution, excit_distribution
101     use proc_pointers, only: update_dmcmc_energy_ptr, update_dmcmc_stag_mag_ptr
102     use proc_pointers, only: update_dmcmc_energy_squared_ptr, update_dmcmc_correlation_ptr
103
104     integer, intent(in) :: idet, iteration
105     type(excit) :: excitation
106     real(p) :: unweighted_walker_pop
107
108     ! Get excitation.
109     excitation = get_excitation(walker_dets(:basis_length,idet), &
110                               walker_dets((1+basis_length):total_basis_length,idet))
111
112     ! When performing importance sampling the result is that certain excitation
113     ! levels have smaller psips populations than the true density matrix by some
114     ! factor. In these cases, we want to multiply the psip population by this factor
115     ! to calculate the contribution from these excitation levels correctly.
116
117     ! In the case of no importance sampling, unweighted_walker_pop = walker_population(1,idet)
118     ! and so this change can be ignored.
119     unweighted_walker_pop = walker_population(1,idet)*dmcmc_accumulated_probs(excitation%nexcit)
120
121     ! If diagonal element, add to the trace.
122     if (excitation%nexcit == 0) trace = trace + walker_population(1,idet)
123     ! See which estimators are to be calculated, and call the corresponding procedures.
124     ! Energy
125     if (doing_dmcmc_calc(dmcmc_energy)) call update_dmcmc_energy_ptr&
126         &(idet, excitation, unweighted_walker_pop)
127     ! Energy squared
128     if (doing_dmcmc_calc(dmcmc_energy_squared)) call update_dmcmc_energy_squared_ptr&
129         &(idet, excitation, unweighted_walker_pop)

```

```

130      ! Spin-spin correlation function
131      if (doing_dmcmc_calc(dmcmc_correlation)) call update_dmcmc_correlation_ptr&
132          &(idet, excitation, unweighted_walker_pop)
133      ! Staggered magnetisation
134      if (doing_dmcmc_calc(dmcmc_staggered_magnetisation)) call update_dmcmc_stag_mag_ptr&
135          &(idet, excitation, unweighted_walker_pop)
136      ! Reduced density matrix
137      if (doing_reduced_dm .and. iteration > start_averaging) &
138          call update_reduced_density_matrix_heisenberg(idet, unweighted_walker_pop)
139      ! Excitation distribution
140      if (calculate_excit_distribution) excit_distribution(excitation%nexcit) = &
141          excit_distribution(excitation%nexcit) + abs(walker_population(1,idet))
142      ! Excitation_distribtuion for calculating importance sampling weights
143      if (dmcmc_find_weights .and. iteration > start_averaging) excit_distribution(excitation%nexcit) = &
144          excit_distribution(excitation%nexcit) + abs(walker_population(1,idet))
145
146  end subroutine call_dmcmc_estimators
147
148  subroutine dmcmc_energy_heisenberg(idet, excitation, walker_pop)
149
150      ! For the Heisenberg model only.
151      ! Add the contribution from the current density matrix element
152      ! to the thermal energy estimate.
153
154      ! In:
155      ! idet: Current position in the main bitstring (density matrix) list.
156      ! excitation: excit type variable which stores information on
157      ! the excitation between the two bitstring ends, corresponding
158      ! to the two labels for the density matrix element.
159      ! walker_pop: number of particles on the current density matrix
160      ! element.
161
162      use basis, only: basis_length, total_basis_length
163      use basis, only: bit_lookup
164      use excitations, only: excit
165      use fciqmc_data, only: walker_dets
166      use fciqmc_data, only: walker_data, H00
167      use fciqmc_data, only: estimator_numerators, energy_index
168      use hubbard_real, only: connected_orbs
169      use system, only: J_coupling
170
171      integer, intent(in) :: idet
172      type(excit), intent(in) :: excitation
173      real(p) , intent(in) :: walker_pop

```

```

174     integer :: bit_element, bit_position
175
176     ! If no excitation, we have a diagonal element, so add elements
177     ! which involve the diagonal element of the Hamiltonian.
178     if (excitation%nexcit == 0) then
179         estimator_numerators(energy_index) = estimator_numerators(energy_index) + &
180             (walker_data(1,idet)+H00)*walker_pop
181     else if (excitation%nexcit == 1) then
182         ! If not a diagonal element, but only a single excitation, then the corresponding
183         ! Hamiltonian element may be non-zero. Calculate if the flipped spins are
184         ! neighbours on the lattice, and if so, add the contirbution from this site.
185         bit_position = bit_lookup(1,excitation%from_orb(1))
186         bit_element = bit_lookup(2,excitation%from_orb(1))
187         if (btest(connected_orbs(bit_element, excitation%to_orb(1)), bit_position)) &
188             estimator_numerators(energy_index) = estimator_numerators(energy_index) - &
189                 (2.0*J_coupling*walker_pop)
190     end if
191
192 end subroutine dmqmc_energy_heisenberg
193
194 subroutine dmqmc_energy_squared_heisenberg(idet, excitation, walker_pop)
195
196     ! For the Heisenberg model only.
197     ! Add the contribution from the current density matrix element
198     ! to the thermal energy squared estimate.
199
200     ! In:
201     ! idet: Current position in the main bitstring (density matrix) list.
202     ! excitation: excit type variable which stores information on
203     ! the excitation between the two bitstring ends, corresponding
204     ! to the two labels for the density matrix element.
205     ! walker_pop: number of particles on the current density matrix
206     ! element.
207
208     use basis, only: basis_length, total_basis_length
209     use basis, only: bit_lookup
210     use excitations, only: excit
211     use fciqmc_data, only: walker_dets
212     use fciqmc_data, only: walker_data, H00
213     use fciqmc_data, only: estimator_numerators, energy_squared_index
214     use hubbard_real, only: connected_orbs, next_nearest_orbs
215     use system, only: J_coupling, nbonds
216
217     integer, intent(in) :: idet

```

```

218     type(excit), intent(in) :: excitat
219     real(p) , intent(in) :: walker_pop
220     integer :: bit_element1, bit_position1, bit_element2, bit_position2
221     real(p) :: sum_H1_H2, J_coupling_squared
222
223     sum_H1_H2 = 0
224     J_coupling_squared = J_coupling**2
225
226     if (excitation%nexcit == 0) then
227         ! If there are 0 excitations then either nothing happens twice, or we
228         ! flip the same pair of spins twice. The Hamiltonian element for doing nothing
229         ! is just the diagonal element. For each possible pairs of spins which can be
230         ! flipped, there is a matrix element of -2*J_coupling, so we just need to count
231         ! the number of such pairs, which can be found simply from the diagonal element.
232
233         sum_H1_H2 = (walker_data(1,idet)+H00)**2
234         sum_H1_H2 = sum_H1_H2 + 2.0*J_coupling_squared*nbonds + 2.0*J_coupling*(walker_data(1,idet)+H00)
235
236     else if (excitation%nexcit == 1) then
237         ! If there is only one excitation (2 spins flipped) then the contribution to H^2
238         ! depend on the positions of the spins relative to one another.
239         ! If the the spins are nearest neighbors then we could either do nothing and then
240         ! flip the pair, or flip the pair and then do nothing.
241         ! If next nearest neighbors and there is only one two-bond path to get from one
242         ! spin to the other, we first flip the pair on the first bond, then flip the pair
243         ! on the second bond. This flipping can only be done in exactly one order, not both -
244         ! the two spins which change are opposite, so the middle spin will initially only
245         ! be the same as one or the other spin. This is nice, because we don't have check
246         ! which way up the intermediate spin is - there will always be one order which contributes.
247         ! If there are two such paths, then this could happen by either paths, but again, the two
248         ! intermediate spins will only allow one order of spin flipping for each path, no
249         ! matter which way up they are, so we only need to check if there are two possible paths.
250
251         if (next_nearest_orbs(excitation%from_orb(1),excitation%to_orb(1)) /= 0) then
252             ! Contribution for next-nearest neighbors.
253             sum_H1_H2 = 4.0*J_coupling_squared*next_nearest_orbs(excitation%from_orb(1),excitation%to_orb(1))
254         end if
255         ! Contributions for nearest neighbors.
256         ! Note, for certain lattices, such as the triangular lattice, two spins can be both
257         ! nearest neighbors *and* next-nearest neighbors. Therefore, it is necessary in general
258         ! to check for both situations.
259         bit_position1 = bit_lookup(1,excitation%from_orb(1))
260         bit_element1 = bit_lookup(2,excitation%from_orb(1))
261         if (btest(connected_orbs(bit_element1, excitat

```

```

262         sum_H1_H2 = sum_H1_H2 - 4.0*J_coupling*(walker_data(1,idet)+H00)
263
264     else if (excitation%nexcit == 2) then
265         ! If there are two excitations (4 spins flipped) then, once again, the contribution
266         ! to the thermal energy squared will depend on the positions of the spins. If there
267         ! are two pairs of spins flipped which are separated then there is one way
268         ! for this to happen - by flipping one pair, and then the other (this also requires
269         ! that the two spins within each neighboring pair are opposite, as ever for the
270         ! Heisenberg model). These two flips can happen in either order.
271         ! In some cases the spins may be such that we may pair the spins in more than one way.
272         ! For example, if the four spins are in a square shape, or for a 4-by-4 Heisenberg
273         ! model, the spins could be connected across the whole lattice, forming a ring due
274         ! to the periodic boundaries. In these cases it may be possible to perform the spin
275         ! flips by pairing them in either of two ways. To account for this possibility we
276         ! have to try and pair the spins in both ways, so we always check both if statements
277         ! below. Again, once these pairings have been chosen, the flips can be performed in
278         ! either order.
279
280         bit_position1 = bit_lookup(1,excitation%from_orb(1))
281         bit_element1 = bit_lookup(2,excitation%from_orb(1))
282         bit_position2 = bit_lookup(1,excitation%from_orb(2))
283         bit_element2 = bit_lookup(2,excitation%from_orb(2))
284         if (btest(connected_orbs(bit_element1, excitation%to_orb(1)), bit_position1) .and. &
285             btest(connected_orbs(bit_element2, excitation%to_orb(2)), bit_position2)) &
286             sum_H1_H2 = 8.0*J_coupling_squared
287         if (btest(connected_orbs(bit_element1, excitation%to_orb(2)), bit_position1) .and. &
288             btest(connected_orbs(bit_element2, excitation%to_orb(1)), bit_position2)) &
289             sum_H1_H2 = sum_H1_H2 + 8.0*J_coupling_squared
290
291     end if
292
293     estimator_numerators(energy_squared_index) = estimator_numerators(energy_squared_index) + &
294         sum_H1_H2*walker_pop
295
296 end subroutine dmqmc_energy_squared_heisenberg
297
298 subroutine dmqmc_energy_hub_real(idet, excitation, walker_pop)
299
300     ! For the Heisenberg model only.
301     ! Add the contribution from the current density matrix element
302     ! to the thermal energy estimate.
303
304     ! In:
305     ! idet: Current position in the main bitstring (density matrix) list.

```

```

306      ! excitation: excit type variable which stores information on
307      ! the excitation between the two bitstring ends, corresponding
308      ! to the two labels for the density matrix element.
309      ! walker_pop: number of particles on the current density matrix
310      ! element.
311
312      use basis, only: basis_length, total_basis_length
313      use excitations, only: excit
314      use fciqmc_data, only: walker_dets
315      use fciqmc_data, only: walker_data, H00
316      use fciqmc_data, only: estimator_numerators, energy_index
317      use hamiltonian, only: slater_condon1_hub_real
318
319      integer, intent(in) :: idet
320      type(excit), intent(in) :: excitation
321      real(p) , intent(in) :: walker_pop
322      real(p) :: hmatel
323
324      ! If no excitation, we have a diagonal element, so add elements
325      ! which involve the diagonal element of the Hamiltonian.
326      if (excitation%nexcit == 0) then
327          estimator_numerators(energy_index) = estimator_numerators(energy_index) + &
328              (walker_data(1,idet)+H00)*walker_pop
329      else if (excitation%nexcit == 1) then
330          ! If not a diagonal element, but only a single excitation, then the corresponding
331          ! Hamiltonian element may be non-zero. Calculate if the flipped spins are
332          ! neighbours on the lattice, and if so, add the contirbution from this site.
333          hmatel = slater_condon1_hub_real(excitation%from_orb(1), excitation%to_orb(1), excitation%perm)
334          estimator_numerators(energy_index) = estimator_numerators(energy_index) + &
335              (hmatel*walker_pop)
336      end if
337
338      end subroutine dmqmc_energy_hub_real
339
340      subroutine dmqmc_correlation_function_heisenberg(idet, excitation, walker_pop)
341
342          ! For the Heisenberg model only.
343          ! Add the contribution from the current density matrix element
344          ! to the thermal spin correlation function estimator.
345
346          ! In:
347          ! idet: Current position in the main bitstring (density matrix) list.
348          ! excitation: excit type variable which stores information on
349          ! the excitation between the two bitstring ends, corresponding

```

```

350      ! to the two labels for the density matrix element.
351      ! walker_pop: number of particles on the current density matrix
352      ! element.
353
354      use basis, only: basis_length, total_basis_length
355      use basis, only: bit_lookup
356      use bit_utils, only: count_set_bits
357      use excitations, only: excit
358      use fciqmc_data, only: walker_dets
359      use fciqmc_data, only: walker_data, H00, correlation_mask
360      use fciqmc_data, only: estimator_numerators, correlation_index
361      use hubbard_real, only: connected_orbs
362      use system, only: J_coupling
363
364      integer, intent(in) :: idet
365      type(excit), intent(in) :: excitation
366      real(p) , intent(in) :: walker_pop
367      integer(i0) :: f(basis_length)
368      integer :: bit_element1, bit_position1, bit_element2, bit_position2
369      integer :: sign_factor
370
371      ! If no excitation, we want the diagonal element of the correlation function operator.
372      if (excitation%nexcit == 0) then
373          ! If the two spins i and j are the same, the matrix element is +1/4.
374          ! If they are different, the matrix element is -1/4.
375          ! So we want sign_factor to be +1 in the former case, and -1 in the latter case.
376          ! f as calculated below will have 0's at sites other than i and j, and the same values
377          ! as walker_dets at i and j. Hence, if f has two 1's or no 1's, we want
378          ! sign_factor = +1. Else if we have one 1, we want sign_factor = -1.
379          f = iand(walker_dets(:basis_length,idet), correlation_mask)
380          ! Count if we have zero, one or two 1's.
381          sign_factor = sum(count_set_bits(f))
382          ! The operation below will map 0 and 2 to +1, and will map 1 to -1, as is easily checked.
383          sign_factor = (mod(sign_factor+1,2)*2)-1
384          ! Hence sign_factor can be used to find the matrix element, as used below.
385          estimator_numerators(correlation_index) = estimator_numerators(correlation_index) + &
386              (sign_factor*(walker_pop/4))
387      else if (excitation%nexcit == 1) then
388          ! If not a diagonal element, but only a single excitation, then the corresponding
389          ! matrix element will be 1/2 if and only if the two sites which are flipped are
390          ! sites i and j, else it will be 0. We assume that excitations will only be set if
391          ! i and j are opposite (else they could not be flipped, for ms=0).
392          bit_position1 = bit_lookup(1,excitation%from_orb(1))
393          bit_element1 = bit_lookup(2,excitation%from_orb(1))

```

```

394         bit_position2 = bit_lookup(1,excitation%to_orb(1))
395         bit_element2 = bit_lookup(2,excitation%to_orb(1))
396         if (btest(correlation_mask(bit_element1), bit_position1).and.btest(correlation_mask(bit_element2),
397             bit_position2)) &
398             estimator_numerators(correlation_index) = estimator_numerators(correlation_index) + &
399                 (walker_pop/2)
400     end if
401 end subroutine dmqmc_correlation_function_heisenberg
402
403 subroutine dmqmc_stag_mag_heisenberg(idet, excitation, walker_pop)
404
405     ! For the Heisenberg model only.
406     ! Add the contribution from the current density matrix element
407     ! to the thermal staggered magnetisation estimate.
408
409     ! In:
410     ! idet: Current position in the main bitstring (density matrix) list.
411     ! excitation: excit type variable which stores information on
412     ! the excitation between the two bitstring ends, corresponding
413     ! to the two labels for the density matrix element.
414     ! walker_pop: number of particles on the current density matrix
415     ! element.
416
417     use basis, only: basis_length, total_basis_length, bit_lookup
418     use bit_utils, only: count_set_bits
419     use determinants, only: lattice_mask
420     use excitations, only: excit
421     use fciqmc_data, only: walker_dets
422     use fciqmc_data, only: estimator_numerators, staggered_mag_index
423     use system, only: nel, nsites
424
425     integer, intent(in) :: idet
426     type(excit), intent(in) :: excitation
427     real(p) , intent(in) :: walker_pop
428     integer :: bit_element1, bit_position1, bit_element2, bit_position2
429     integer(i0) :: f(basis_length)
430     integer :: n_up_plus
431     integer :: total_sum
432
433     total_sum = 0
434
435     ! This is for the staggered magnetisation squared. This is given by
436     !  $M^2 = M_x^2 + M_y^2 + M_z^2$ 

```

```

437      ! where  $M_x = \sum_i (-1)^i S_i^x$ , etc... and  $S_i^x$  is the spin operator
438      ! at site  $i$ , in the  $x$  direction.
439      if (excitation%nextcit == 0) then
440          ! Need to calculate the number of spins up on sublattice 1:
441          !  $N_u(+)$  - Number up on + sublattice (where  $(-1)^i$  above is +1)
442          ! Note the number of down spins on a sublattice is easily obtained from
443          !  $N_u(+)$  since there are  $N/2$  spins on each - and the number of spins up on
444          ! a different sublattice is easily obtained since there are  $n_{el}$  spins
445          ! up in total. Hence the matrix element will be written only in terms
446          ! of the number of up spins on sublattice 1, to save computation.
447          f = iand(walker_dets(:basis_length,idet), lattice_mask)
448          n_up_plus = sum(count_set_bits(f))
449          ! Below, the term in brackets and middle term come from the  $z$  component (the
450          !  $z$  operator is diagonal) and one  $nsites/4$  factor comes from the  $x$  operator,
451          ! the other  $nsites/4$  factor from the  $y$  operator.
452          total_sum = (2*n_up_plus-nel)**2 + (nsites/2)
453      else if (excitation%nextcit == 1) then
454          ! Off-diagonal elements from the  $y$  and  $z$  operators. For the pair of spins
455          ! that are flipped, if they are on the same sublattice, we get a factor of
456          ! 1, or if on different sublattices, a factor of -1.
457          bit_position1 = bit_lookup(1,excitation%from_orb(1))
458          bit_element1 = bit_lookup(2,excitation%from_orb(1))
459          bit_position2 = bit_lookup(1,excitation%to_orb(1))
460          bit_element2 = bit_lookup(2,excitation%to_orb(1))
461          if (btest(lattice_mask(bit_element1), bit_position1)) total_sum = total_sum+1
462          if (btest(lattice_mask(bit_element2), bit_position2)) total_sum = total_sum+1
463          ! The operation below will map 0 and 2 to +1, and will map 1 to -1, as is easily checked.
464          ! We want this - if both or no spins on this sublattice, then both on same sublattice
465          ! either way, so plus one. Else they are on different sublattices, so we want a factor
466          ! of -1, as we get.
467          total_sum = (mod(total_sum+1,2)*2)-1
468      end if
469
470      estimator_numerators(staggered_mag_index) = estimator_numerators(staggered_mag_index) + &
471          (real(total_sum)/real(nsites**2))*walker_pop
472
473      end subroutine dmqmc_stag_mag_heisenberg
474
475      subroutine update_reduced_density_matrix_heisenberg(idet, walker_pop)
476
477          ! Add a contribution from the current walker to the reduced density
478          ! matrix estimator, which is produced by 'tracing out' a given set of
479          ! spins.
480

```

```

481      ! Applicable only to the Heisenberg model.
482
483      ! This procedure takes the two determinants bitstrings for the current
484      ! walker and, if the two bitstrings of the B subsystem are identical,
485      ! adds the walker population to the corresponding reduced density matrix
486      ! element.
487
488      ! In:
489      ! idet: Current position in the main bitstring (density matrix) list.
490      ! walker_pop: number of particles on the current density matrix
491      ! element.
492
493      use basis, only: basis_length, total_basis_length
494      use dmqmc_procedures, only: decode_dm_bitstring
495      use fciqmc_data, only: subsystem_B_mask, reduced_density_matrix
496      use fciqmc_data, only: walker_dets, walker_population
497
498      integer, intent(in) :: idet
499      real(p), intent(in) :: walker_pop
500      integer(i0) :: f1(basis_length), f2(basis_length)
501      integer(i0) :: end1, end2
502
503      ! Apply the mask for the B subsystem to set all sites in the A subsystem to 0.
504      f1 = iand(subsystem_B_mask,walker_dets(:basis_length,idet))
505      f2 = iand(subsystem_B_mask,walker_dets(basis_length+1:total_basis_length,idet))
506
507      ! Once this is done, check if the resulting bitstring (which can only possibly
508      ! have 1's in the B subsystem) are identical. If they are, then this psip gives
509      ! a contribution to the reduced density matrix for subsystem A. This is because we
510      ! get the reduced density matrix for A by 'tracing out' over B, which in practice
511      ! means only keeping matrix elements that are on the diagonal for subsystem B.
512      if (sum(abs(f1-f2)) == 0) then
513          ! We need to assign positions in the reduced density matrix for this walker,
514          ! so call a function which maps the possible subsystem A spin configurations
515          ! to indices.
516          call decode_dm_bitstring(walker_dets(:,idet),end1,end2)
517          reduced_density_matrix(end1,end2) = reduced_density_matrix(end1,end2) + walker_pop
518      end if
519
520      end subroutine update_reduced_density_matrix_heisenberg
521
522      subroutine call_rdm_procedures()
523
524          ! Wrapper for calling relevant reduced density matrix procedures

```

```

525     use fciqmc_data, only: reduced_density_matrix, subsystem_A_size
526     use fciqmc_data, only: doing_von_neumann_entropy, doing_concurrence
527     use parallel
528
529     real(p) :: trace_rdm
530     integer :: i, j
531     ! If in paralell then merge the reduced density matrix onto one processor
532 #ifdef PARALLEL
533
534     real(dp) :: dm(2**subsystem_A_size,2**subsystem_A_size)
535     real(dp) :: dm_sum(2**subsystem_A_size,2**subsystem_A_size)
536     integer :: ierr
537
538     dm = reduced_density_matrix
539
540     call mpi_allreduce(dm, dm_sum, size(dm), MPI_REAL8, MPI_SUM, MPI_COMM_WORLD, ierr)
541
542     reduced_density_matrix = dm_sum
543
544 #endif
545
546     trace_rdm = 0.0_p
547
548     if (parent) then
549         ! Force the reduced desnity matrix to be symmetric by averaging the upper and lower triangles
550         do i = 1, ubound(reduced_density_matrix,1)
551             do j = 1, i-1
552                 reduced_density_matrix(i,j) = 0.5_p*(reduced_density_matrix(i,j) +&
553                 reduced_density_matrix(j,i))
554                 reduced_density_matrix(j,i) = reduced_density_matrix(i,j)
555             end do
556             ! Add current contirbution to the trace.
557             trace_rdm = trace_rdm + reduced_density_matrix(i,i)
558         end do
559
560         ! Call the routines to calculate the desired quantities.
561         if (doing_von_neumann_entropy) call calculate_vn_entropy()
562         if (doing_concurrence) call calculate_concurrence()
563
564         write (6,'(1x,a12,1X,f22.12)') "# RDM trace=", trace_rdm
565     end if
566
567
568 end subroutine call_rdm_procedures

```

```

569
570     subroutine calculate_vn_entropy()
571
572         ! Calculate the Von Neumann Entropy. Use lapack to calculate the
573         ! eigenvalues {\lambda_j} of the reduced density matrix.
574         ! Then VN Entropy S = -\sum_j \lambda_j \log_2 \{\lambda_j\}
575
576         ! Need to parallise for large subsystems and introduce test to
577         ! check whether diagonalisation should be performed in serial or
578         ! parallel.
579
580         use checking, only: check_allocate, check_deallocate
581         use fciqmc_data, only: subsystem_A_size, reduced_density_matrix
582
583         integer :: i, j, rdm_size
584         integer :: info, ierr, lwork
585         real(p), allocatable :: work(:)
586         real(p) :: eigv(2*subsystem_A_size)
587         real(p) :: vn_entropy
588
589         rdm_size = 2*subsystem_A_size
590         vn_entropy = 0._p
591
592         ! Find the optimal size of the workspace.
593         allocate(work(1), stat=ierr)
594         call check_allocate('work',1,ierr)
595 #ifdef SINGLE_PRECISION
596         call ssyev('N', 'U', rdm_size, reduced_density_matrix, rdm_size, eigv, work, -1, info)
597 #else
598         call dsyev('N', 'U', rdm_size, reduced_density_matrix, rdm_size, eigv, work, -1, info)
599 #endif
600         lwork = nint(work(1))
601         deallocate(work)
602         call check_deallocate('work',ierr)
603
604         ! Now perform the diagonalisation.
605         allocate(work(lwork), stat=ierr)
606         call check_allocate('work',lwork,ierr)
607
608 #ifdef SINGLE_PRECISION
609         call ssyev('N', 'U', rdm_size, reduced_density_matrix, rdm_size, eigv, work, lwork, info)
610 #else
611         call dsyev('N', 'U', rdm_size, reduced_density_matrix, rdm_size, eigv, work, lwork, info)
612 #endif

```

```

613         do i = 1, ubound(eigv,1)
614             vn_entropy = vn_entropy - eigv(i)*(log(eigv(i))/log(2.0_p))
615         end do
616         write (6,'(1x,a36,1X,f22.12)') "# Unnormalised Von-Neumann Entropy= ", vn_entropy
617
618     end subroutine calculate_vn_entropy
619
620     subroutine calculate_concurrence()
621
622         ! Calculate the concurrence of a qubit. For a reduced density matrix \rho,
623         ! the concurrence, C = max(0, \lambda_1 - \lambda_2 - \lambda_3 - \lambda_4) where
624         ! \lambda_i are the eigenvalues of the matrix, R = \sqrt{\sqrt{\rho}\rho\sqrt{\rho}}.
625         ! Where \rho = {\sigma_y \otimes \sigma_y} \rho^{\ast} {\sigma_y \otimes \sigma_y}.
626         ! \lambda_1 > ... > \lambda_4
627
628         ! This can be simplified to finding the square root of the eigenvalues {\lambda_i} of \rho^{\rho}
629         ! and in the case where \rho is a real, symmetric matrix then we can further simplify the problem
630         ! to finding the eigenvalues of R = \rho \sigma_y \otimes \sigma_y.
631
632         ! Below we have named {\sigma_y \otimes \sigma_y} flip_spin_matrix as in the literature.
633
634         use checking, only: check_allocate, check_deallocate
635         use fciqmc_data, only: subsystem_A_size, reduced_density_matrix, flip_spin_matrix
636         use utils, only: print_matrix
637         integer :: i,j
638         integer :: info, ierr, lwork
639         real(p), allocatable :: work(:)
640         real(p) :: reigv(4), ieigv(4)
641         real(p) :: concurrence
642         real(p) :: rdm_spin_flip(4,4), rdm_spin_flip_tmp(4,4), VL(4,4), VR(4,4)
643
644         ! Make rdm_spin_flip_tmp because sgeev and dgeev delete input matrix
645         rdm_spin_flip_tmp = matmul(reduced_density_matrix, flip_spin_matrix)
646         rdm_spin_flip = rdm_spin_flip_tmp
647
648         ! Find the optimal size of the workspace.
649         allocate(work(1), stat=ierr)
650         call check_allocate('work',1,ierr)
651     #ifdef SINGLE_PRECISION
652         call sgeev('N', 'N', 4, rdm_spin_flip_tmp, 4, reigv, ieigv, VL, 1, VR, 1, work, -1, info)
653     #else
654         call dgeev('N', 'N', 4, rdm_spin_flip_tmp, 4, reigv, ieigv, VL, 1, VR, 1, work, -1, info)
655     #endif
656         lwork = nint(work(1))

```

```

657         deallocate(work)
658         call check_deallocate('work',ierr)
659
660         ! Now perform the diagonalisation.
661         allocate(work(lwork), stat=ierr)
662         call check_allocate('work',lwork,ierr)
663 #ifdef SINGLE_PRECISION
664         call sgeev('N', 'N', 4, rdm_spin_flip, 4, reigv, ieigv, VL, 1, VR, 1, work, lwork, info)
665 #else
666         call dgeev('N', 'N', 4, rdm_spin_flip, 4, reigv, ieigv, VL, 1, VR, 1, work, lwork, info)
667 #endif
668         ! Calculate the concurrence. Take abs of eigenvalues so that this is equivalent to squaring
669         ! and then square-rooting
670         concurrence = 2._p*maxval(abs(reigv)) - sum(abs(reigv))
671         concurrence = max(0._p, concurrence)
672         write (6,'(1x,a28,1X,f22.12)') "# Unnormalised concurrence= ", concurrence
673
674     end subroutine calculate_concurrence
675
676 end module dmqmc_estimators

```
